

The Entity-Relationship Model

COMPSCI 2DB3: Databases

Jelle Hellings Holly Koponen

Department of Computing and Software
McMaster University



Winter 2024

The relational data model

Data model

The rules by which real-world data can be represented and structured.

The relational data model

All data is modeled as a collection of *tables*.

Terminology: Schema and Instance

Schema Describe how stored data is structured.

- ▶ Conceptual schema: in terms of the *data model*.
- ▶ Physical schema: details in terms of *layout on disk*: files, file structure, auxiliary data structures (indices),

Instance The actual data (adhering to some schema).

Terminology: Schema and Instance

Schema Describe how stored data is structured.

- ▶ Conceptual schema: in terms of the *data model*.
- ▶ Physical schema: details in terms of *layout on disk*: files, file structure, auxiliary data structures (indices),

Instance The actual data (adhering to some schema).

Example: The relational data model

cgrades	(<u>id</u> ,	name,	program,	age,	score)
----------------	---------------	-------	----------	------	--------

The relational schema describing a table

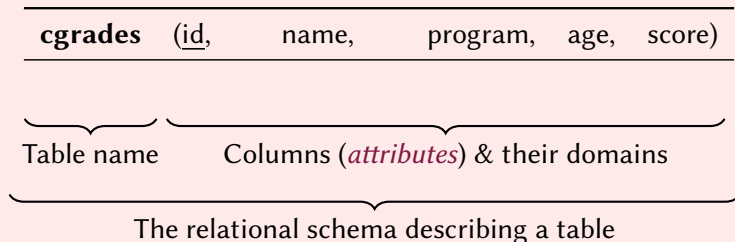
Terminology: Schema and Instance

Schema Describe how stored data is structured.

- ▶ Conceptual schema: in terms of the *data model*.
- ▶ Physical schema: details in terms of *layout on disk*: files, file structure, auxiliary data structures (indices),

Instance The actual data (adhering to some schema).

Example: The relational data model



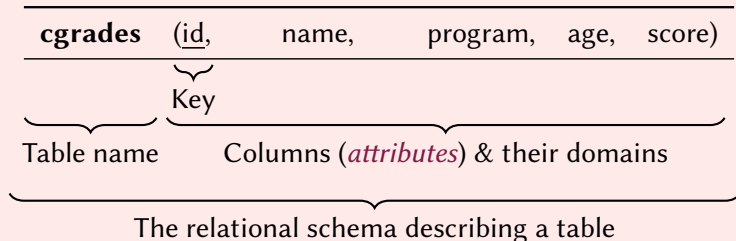
Terminology: Schema and Instance

Schema Describe how stored data is structured.

- ▶ Conceptual schema: in terms of the *data model*.
- ▶ Physical schema: details in terms of *layout on disk*: files, file structure, auxiliary data structures (indices),

Instance The actual data (adhering to some schema).

Example: The relational data model



Terminology: Schema and Instance

Schema Describe how stored data is structured.

- ▶ Conceptual schema: in terms of the *data model*.
- ▶ Physical schema: details in terms of *layout on disk*: files, file structure, auxiliary data structures (indices),

Instance The actual data (adhering to some schema).

Example: The relational data model

cgrades	(<u>id</u> ,	name,	program,	age,	score)	
	53666	Jones	cs	18	3.4	} Instance
	53688	Smith	ee	18	3.2	
	53650	Smith	math	19	3.8	
	53831	Madayan	music	11	1.8	
	53832	Guldu	music	12	2.0	

How to create a schema?

Requirements Analysis

- ▶ What data can this application collect?
- ▶ What data does this application need?
- ▶ What does the application do with the data?

How to create a schema?

Requirements Analysis

- ▶ What data can this application collect?
- ▶ What data does this application need?
- ▶ What does the application do with the data?

Is a database management system the right fit?

E.g., big data analytics, machine learning, visualization, ...

How to create a schema?

Requirements Analysis

- ▶ What data can this application collect?
- ▶ What data does this application need?
- ▶ What does the application do with the data?

Is a database management system the right fit?

E.g., big data analytics, machine learning, visualization, ...

What about agile development?

It is complex, costly, and sometimes impossible to refactor away *wrong data choices*. E.g.,

- ▶ How do you add missing data after the fact?
- ▶ How do you restructure data to improve performance?

Agile development *requires* expertise to make the right choices at the right time.

Requirements Analysis: the entity-relationship data model

Relational data: Tables represent everything

Faculty(fid, name, location, mail),

Course(cid, name, year, credits).

Teaches(fid, cid).

Requirements Analysis: the entity-relationship data model

Relational data: Tables represent everything

<i>Faculty</i> (<u>fid</u> , name, location, mail),	}	<i>Entities</i>
<i>Course</i> (<u>cid</u> , name, year, credits).		
<i>Teaches</i> (<u>fid</u> , <u>cid</u>).	}	<i>Relationship</i>

Requirements Analysis: the entity-relationship data model

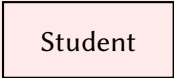
Relational data: Tables represent everything

<i>Faculty</i> (<u>fid</u> , name, location, mail),	}	<i>Entities</i>
<i>Course</i> (<u>cid</u> , name, year, credits).		
<i>Teaches</i> (<u>fid</u> , <u>cid</u>).	}	<i>Relationship</i>

Entity-Relationship data model

- ▶ A *semantic* data model: Specify *meaning* and not *representation*.
- ▶ Can model high-level *concepts*: entities, relationships, attributes,
- ▶ Can easily be translated to well-designed tables (*representation*).
- ▶ Resulting ER-diagrams are easier to document and discuss with *non-specialists*.

A step-by-step example: Students enrolled in courses




Student

A rectangular box with a black border, containing the word "Student" in a black serif font.


The *entity set* student.

A step-by-step example: Students enrolled in courses




Students have *attributes*: student id, name, and age.

A step-by-step example: Students enrolled in courses




Student id is a *key*: a unique identifier.

A step-by-step example: Students enrolled in courses





Age is a *derived* attribute: birthdate is easier to maintain!

A step-by-step example: Students enrolled in courses



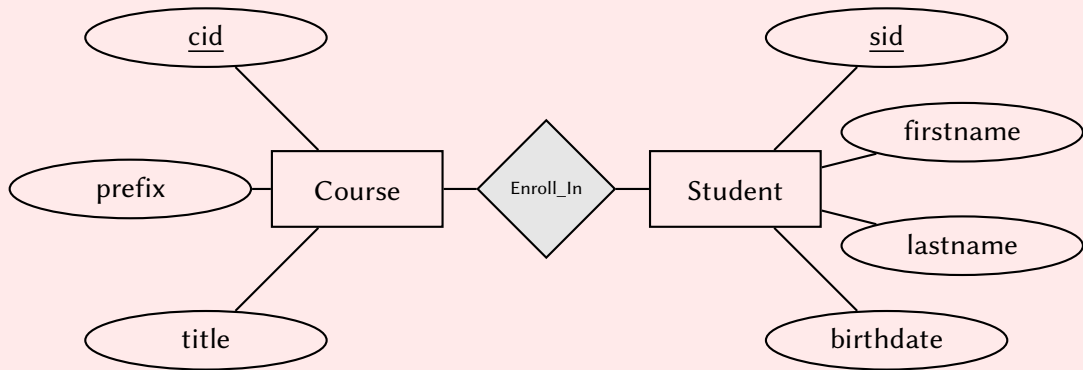
Names are *complex*: does the application need first and last names?

A step-by-step example: Students enrolled in courses



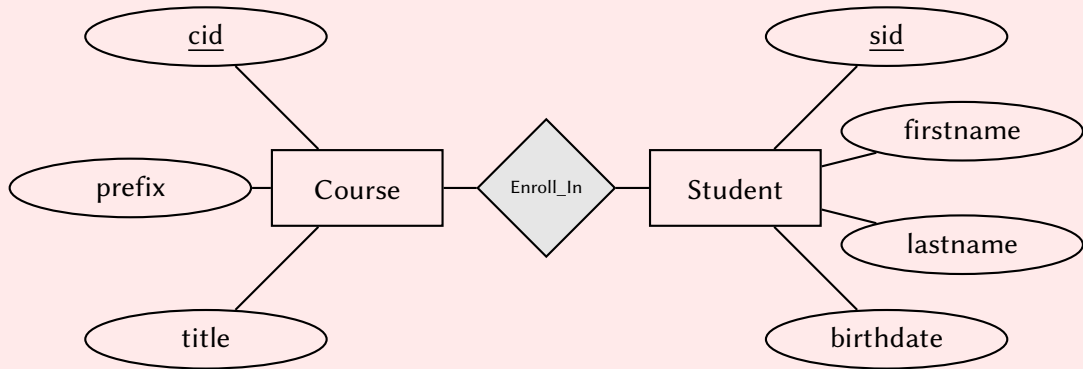
Another *entity*: a course.

A step-by-step example: Students enrolled in courses



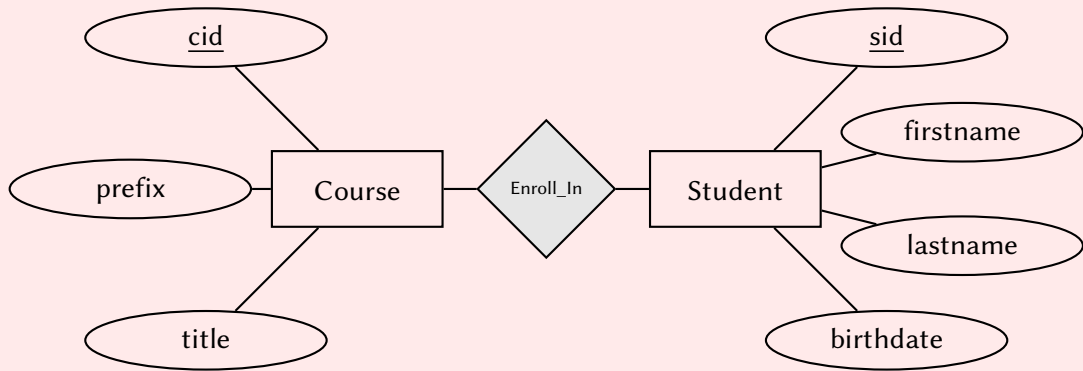
Each student can enroll in classes: a *relationship set*.

A step-by-step example: Students enrolled in courses



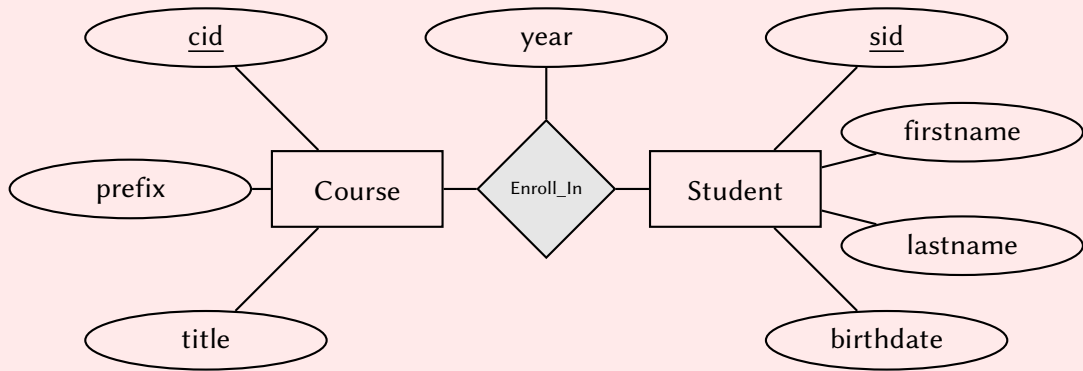
This is a *many-to-many relationship set*.

A step-by-step example: Students enrolled in courses



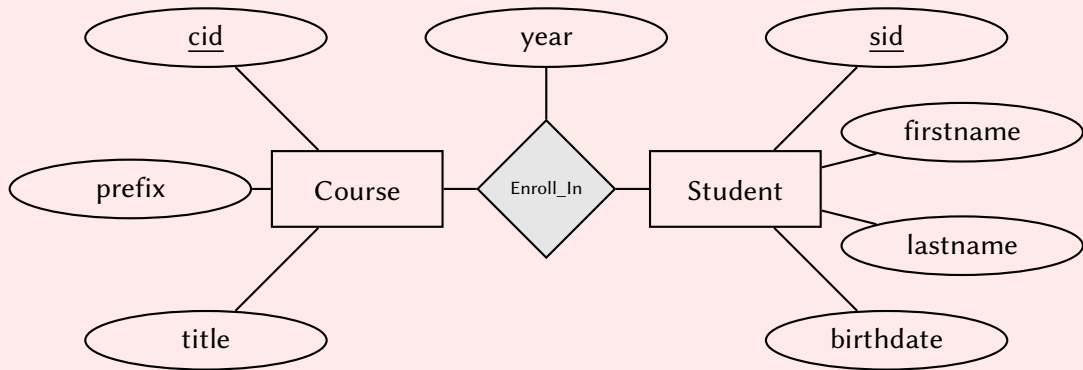
Relationships are *unique*: A student cannot take a single course multiple times.

A step-by-step example: Students enrolled in courses



Relationship sets can have attributes: the *year* in which the course was taken.

A step-by-step example: Students enrolled in courses



Even with relationship attributes: Students can take a course only once!

Intermission: Names...

Question: Did you ever have problems with your name?

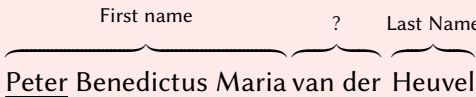
Vote at <https://strawpoll.com/cp6fysdxr>.

Or: go to <https://strawpoll.live> and use the code **112128**.

Intermission: Names—an example


First name ? Last Name


The Netherlands (Dutch) Peter Benedictus Maria van der Heuvel.



Some useful background reading: [Falsehoods Programmers Believe About Names](#).


Intermission: Names—an example


The Netherlands (Dutch)  Peter Benedictus Maria van der Heuvel.

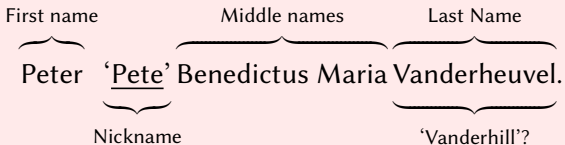
Belgium (Dutch)  Peter ~~Benedictus Maria~~ Van der Heuvel.

Some useful background reading: [Falsehoods Programmers Believe About Names](#).

Intermission: Names—an example

The Netherlands (Dutch)  Peter Benedictus Maria van der Heuvel.

Belgium (Dutch)  Peter ~~Benedictus~~ ~~Maria~~ Van der Heuvel.
(rare in Belgium)

USA (English)  Peter 'Pete' Benedictus Maria Vanderheuve.
Nickname 'Vanderhill'?

Some useful background reading: [Falsehoods Programmers Believe About Names](#).


Intermission: What to do with names?

Ask yourself: Why do you need the name?

1. Do not put up limitations, unless you have a specific goal.
2. Be clear to users so they can provide correct info. E.g.,
 - Informal Ask *what they want* to be called (e.g., Hi!).
 - Formal Ask *what they want* to be called formally (e.g., letters).
 - Billing Ask the name *as used by their bank*.
 - Shipping Ask the name *as used by the mail carrier*.
 - Travel Ask the name *as printed in their travel documents*.
3. Have *sensible defaults*: e.g., prefill billing names with account information.


There is no standard solution that works in all cases!

The entity-relationship model: Entities




- ▶ An *entity* is an object in the real world.
- ▶ Entities are described via a set of *attributes*.
- ▶ Attributes have a *domain* (e.g., text, number, date).

The entity-relationship model: Entities




- ▶ An *entity* is an object in the real world.
- ▶ Entities are described via a set of *attributes*.
- ▶ Attributes have a *domain* (e.g., text, number, date).
- ▶ *Key*: set of attributes that identify an entity.
- ▶ All keys are *candidate keys*.
- ▶ One is chosen as the *primary key*.

The entity-relationship model: Entities



- ▶ An *entity* is an object in the real world.
- ▶ Entities are described via a set of *attributes*.
- ▶ Attributes have a *domain* (e.g., text, number, date).
- ▶ *Key*: set of attributes that identify an entity.
- ▶ All keys are *candidate keys*.
- ▶ One is chosen as the *primary key*.
- ▶ An *entity set*: collection of similar entities.
- ▶ An *instance*: snapshot of the entities (“data”).


The entity-relationship model: Entities



Instance


<u>sid</u>	name	birthdate
53666	Jones	May 18, 1999
53688	Smith	June 2, 1994
53650	Smith	December 21, 1995
53831	Madayan	February 7, 2000
53832	Guldu	April 1, 1991

The entity-relationship model: Relationships




- ▶ A *relationship* relates two or more entities.
- ▶ Standard relationship: *many-to-many*:
Students can enroll in *many* courses, courses can have *many* enrolled students.
- ▶ Relationships can have additional *attributes*.

The entity-relationship model: Relationships



- ▶ A *relationship* relates two or more entities.
- ▶ Standard relationship: *many-to-many*:
Students can enroll in *many* courses, courses can have *many* enrolled students.
- ▶ Relationships can have additional *attributes*.
- ▶ A *relationship set*: collection of similar relationships between entities.
- ▶ A *instance*: snapshot of the relations (“data”).

The entity-relationship model: Relationships



Instance


<i>Course</i>	<i>Student</i>	year	<i>Course</i>	<i>Student</i>	year
<i>Course_1</i>	<i>Student_1</i>	2019	<i>Course_2</i>	<i>Student_1</i>	2019
<i>Course_1</i>	<i>Student_2</i>	2020	<i>Course_2</i>	<i>Student_3</i>	2021

More on relationships

Relationships are at the core of the entity-relationship model.

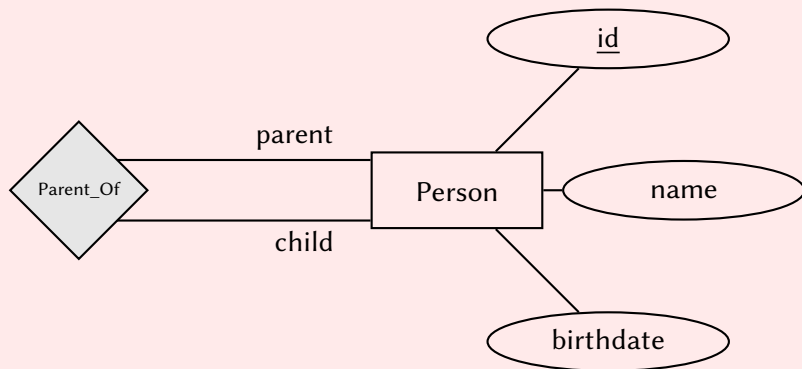
- ▶ Self-referential relationship sets.
- ▶ n -ary relationship sets.
- ▶ Key constraints.

Self-relationships



- Entities in an entity set can be related to each others.
E.g., *Persons* can be *parents*, *friends*, *partners*, *colleagues*.


Self-relationships



- ▶ Entities in an entity set can be related to each others.
E.g., *Persons* can be *parents*, *friends*, *partners*, *colleagues*.
- ▶ Typically, one names the *roles* in such relationships sets.

Ternary relationships

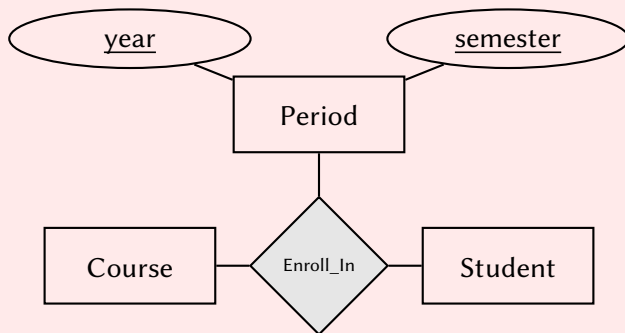
What if a student can enroll several times?



Pairs (*student*, *course*) must be unique.

Ternary relationships

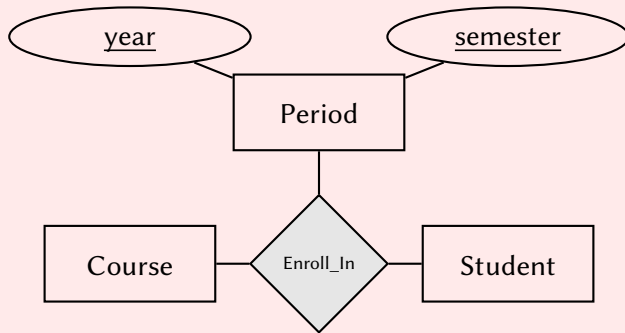
What if a student can enroll several times?



Pairs (*student*, *course*, *period*) must be unique.

Ternary relationships

What if a student can enroll several times?



Ternary relationships are rare: this can be modeled in many other ways.

Types of binary relations (2DM3)

A binary relation $R \subseteq A \times B$ is

Many-to-Many if there are no restrictions on the relation.


One-to-Many if each $a \in A$ is related to *at-most-one* $b \in B$.

- ▶ If A has *total participation*: R is a *function* of A ,
each $a \in A$ is related to *exactly-one* $b \in B$.

One-to-One each A is related to *at-most-one* B and vice-versa.


- ▶ If A has *total participation*: R is an *injection*,
each $a \in A$ is related to exactly-one-and-unique $b \in B$.
- ▶ If B has *total participation*: R is an *surjection*,
each $b \in B$ is related to exactly-one-and-unique $a \in A$.
- ▶ If A, B have *total participation*: R is a *bijection*,
there is a one-to-one mapping between all $a \in A$ s and $b \in B$ s.

Constraints on relationships




A person *likes* many dishes, but has *one* favorite dish.

Constraints on relationships




This says: a person has *many* favorite dishes!

Constraints on relationships



Key constraint: A person has *at-most-one* favorite dish.


Constraints on relationships



Key constraint: A person has *at-most-one* favorite dish.


$E \longrightarrow R$: entity E partakes at-most-once in a R -relationship (partial).

Constraints on relationships



Key constraint: A person has *exactly-one* favorite dish.


Constraints on relationships



Key constraint: A person has *exactly-one* favorite dish.


$E \longrightarrow R$: entity E partakes exactly-once in a R -relationship (total).

Constraints on relationships



Participation constraint: Every dish is *liked*.


Constraints on relationships



Participation constraint: Every dish is *liked*.

E — R : entity E partakes at-least-once in a R -relationship (total).


Constraints on relationships



Remarks on notation

- ▶ We follow the notation of the textbook.
- ▶ Many other sources use different notations (e.g., arrows reversed).
- ▶ We use dedicated notation for exactly-once participation ($\text{---}\rightarrow$).

Key constraints




<i>Entities A</i>
aid
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
bid
1
2
3
4

many-to-many

Key constraints




<i>Entities A</i>
<u>aid</u>
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
<u>bid</u>
1
2
3
4

one-to-many
(partial)

Key constraints




<i>Entities A</i>
<u>aid</u>
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
<u>bid</u>
1
2
3
4

one-to-many
(total)
(function)

Key constraints




<i>Entities A</i>
aid
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
bid
1
2
3
4

many-to-one
(partial)

Key constraints




<i>Entities A</i>
<u>aid</u>
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
<u>bid</u>
1
2
3
4

many-to-one
(total)

Key constraints




<i>Entities A</i>
<u>aid</u>
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
<u>bid</u>
1
2
3
4

one-to-one
(partial, partial)

Key constraints




<i>Entities A</i>
aid
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
bid
1
2
3
4

one-to-one
(total, partial)
(injection)

Key constraints




<i>Entities A</i>
aid
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
bid
1
2
3
4

one-to-one
(partial, total)
(surjection)

Key constraints




<i>Entities A</i>
aid
α
β
γ
δ

<i>Relationships R</i>	
aid	bid
α	1
α	2
β	1
β	2
γ	3

<i>Entities B</i>
bid
1
2
3
4

one-to-one
(total, total)
(bijection)

Key constraints and ternary relationships–1




- ▶ A student can do *one* Master Project.
- ▶ Each project has a supervising professor and is performed in a degree program.

Question: Where should the key constraints go?

Vote at <https://strawpoll.com/3f6eovfz8>.

Or: go to <https://strawpoll.live> and use the code **265451**.

Key constraints and ternary relationships–1




- ▶ A student can do *one* Master Project.
- ▶ Each project has a supervising professor and is performed in a degree program.

Answer: Constraint on Student

- ▶ Key constraint for the participation of students.
- ▶ Professors can supervise *many* Master Projects.
- ▶ *Many* students can perform a Master Project within a program.

Key constraints and ternary relationships–2




Each passenger (customer) of a flight bought a single ticket via a single airline.

Question: Where should the key constraints go?

Vote at <https://strawpoll.com/gwbwsh9hy>.

Or: go to <https://strawpoll.live> and use the code **241367**.

Key constraints and ternary relationships–2



Each passenger (customer) of a flight bought a single ticket via a single airline.

Answer: This is not a key constraint


- ▶ Customers can be on *many* flights.
- ▶ Airlines can sell tickets for *many* flights.
- ▶ Flights can have *many* customers, each from different airlines.

Key constraints and ternary relationships–2

Each passenger (customer) of a flight bought a single ticket via a single airline.

Answer: This is not a key constraint

We can make Passenger a entity!



The basics: entities, relationships, and constraints

We covered enough to model *most typical* situations.

We have already seen situations we could not model, however!

Advanced modeling features

- ▶ Weak entities.
- ▶ ISA Hierarchies.

The need for weak entities

Definition

A *weak entity* is 'owned by' another identity:

A weak entity can only be uniquely identified in conjunction with owning entities.

The need for weak entities

Definition

A *weak entity* is 'owned by' another identity:

A weak entity can only be uniquely identified in conjunction with owning entities.

Examples


Assignments belong to courses:

- ▶ Many courses have a 1-st assignment.
- ▶ The 1-st assignment of COMPSCI 2DB3: *ER data model*.

Courses belong to a degree program (at McMaster):



- ▶ There are several courses with code 2DB3 (e.g., COMPSCI 2DB3 versus SFWRENG 2DB3).
- ▶ There is only one COMPSCI 2DB3 (this course).

Weak entities in practice




A degree program is an entity with a *unique prefix* (COMPSCI, SFWRENG, ...).

Weak entities in practice




A course has a *unique code* within a degree program (2DB3 within COMPSCI).

Weak entities in practice




Course is a *weak entity* and Part_Of is an *identifying relationship* (bold).

Weak entities in practice




Weak entities must participate in their identifying relationship(s): \longrightarrow .

Weak entities in practice




Attribute code is a *partial key* for the weak entity **Course**.

Weak entities in practice




The primary key for weak entity **Course** is the pair (prefix, code)!

Weak entities in practice




The primary key for weak entity Assignment is the triple (prefix, code, number).

A step-by-step example: Continued



Problem: Student could only enroll once

A step-by-step example: Continued




Problem: Student could only enroll once

Turn relationship Enroll_In into a weak entity Enrollment.

- ▶ Enrollment has partial key year.
- ▶ Enrollment is owned by both a Student and a Course.
- ▶ The primary key of Enrollment is the triple (sid, cid, year).


A step-by-step example: Continued



A word of caution

That something *can* be written, does not mean it *should* be written.
Do *not* use weak entities when simpler constructs suffice!

A step-by-step example: Continued



A word of caution

That something *can* be written, does not mean it *should* be written.


Do *not* use weak entities when simpler constructs suffice!

E.g., do not use weak entities to express normal relationships.

Entity modularity via ISAs: Students, Faculty, Staff


This is a classic example of a *class hierarchy*.

Entity modularity via ISAs: Students, Faculty, Staff




Students, Faculty, and Staff are all *peoples* with names and University IDs.

Entity modularity via ISAs: Students, Faculty, Staff




Each student is part of a *generation* (e.g., Class of 2021).

Entity modularity via ISAs: Students, Faculty, Staff




We say that entity Student is a *specialization* or *subclass* of entity Person.

Entity modularity via ISAs: Students, Faculty, Staff



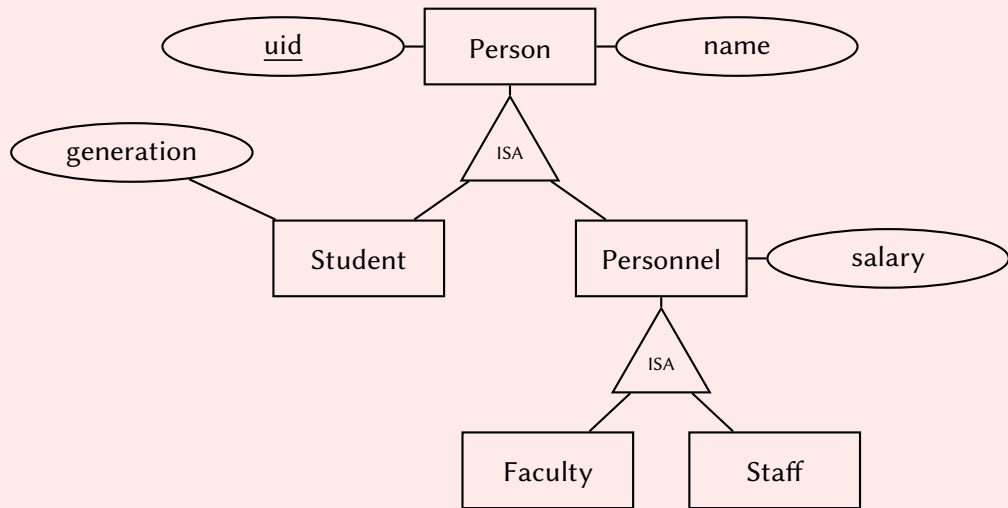
We say that entity Person is a *generalization* or *superclass* of Student.

Entity modularity via ISAs: Students, Faculty, Staff



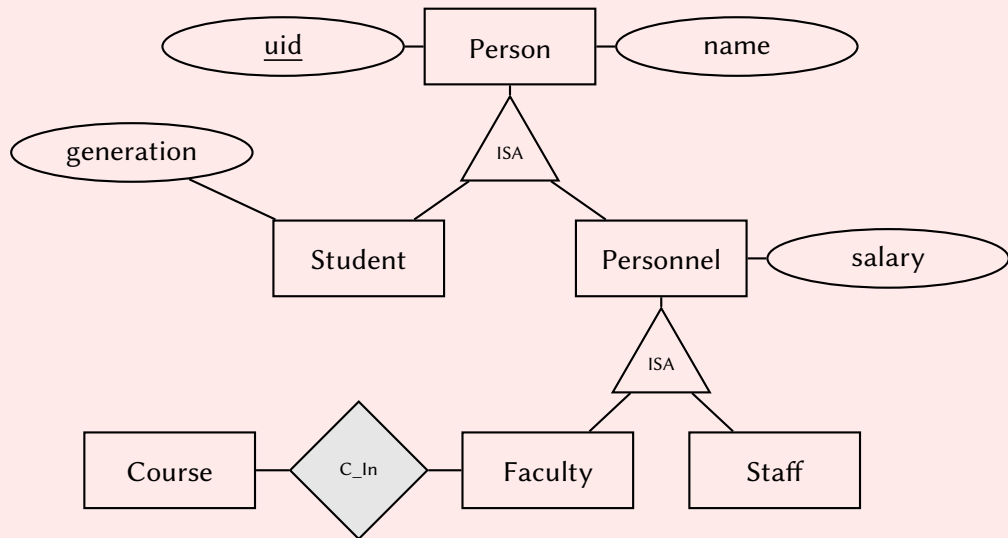
Each faculty and staff member is personnel with a *salary*.

Entity modularity via ISAs: Students, Faculty, Staff



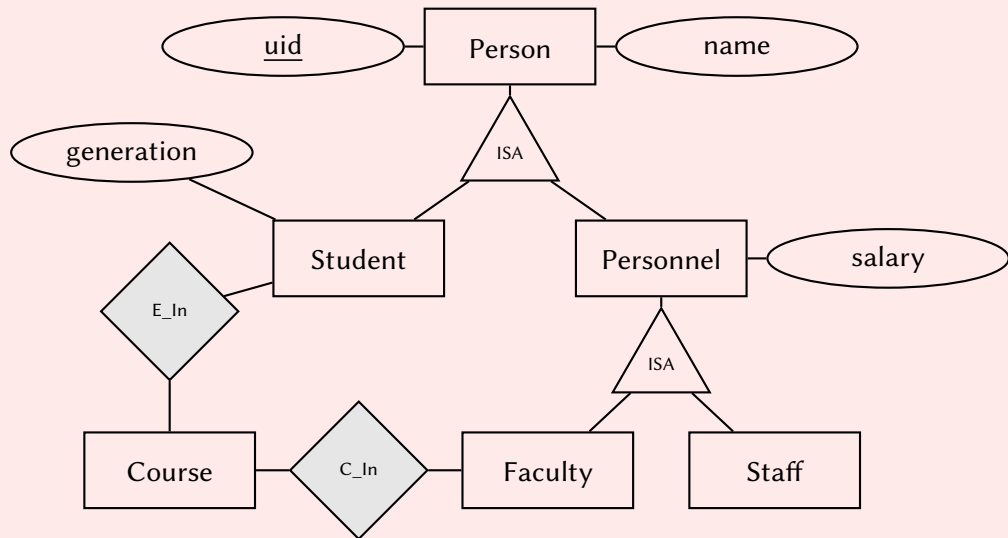
We can further specialize faculty and staff.

Entity modularity via ISAs: Students, Faculty, Staff



Only faculty members are instructors of courses.

Entity modularity via ISAs: Students, Faculty, Staff



While only students can enroll in these courses.

ISA hierarchies require documentation

Every combination of entities in the hierarchy can exist. E.g.,

- ▶ there can be Person entities that are not students, faculty, or staff;
- ▶ there can be Faculty entities that are also Staff and Students,

ISA hierarchies require documentation

Every combination of entities in the hierarchy can exist. E.g.,

- ▶ there can be Person entities that are not students, faculty, or staff;
- ▶ there can be Faculty entities that are also Staff and Students,

Terminology

Overlap constraints Can a single entity belong to *multiple* subclasses?
E.g., can an entity be both Faculty and Student?

Covering constraints Must every superclass entity *also belong* to one of its subclasses?
E.g., must every Person entity be a Student, Faculty, and/or Staff?

ISA hierarchies require documentation

Every combination of entities in the hierarchy can exist. E.g.,

- ▶ there can be Person entities that are not students, faculty, or staff;
- ▶ there can be Faculty entities that are also Staff and Students,

Terminology

Overlap constraints Can a single entity belong to *multiple* subclasses?
E.g., can an entity be both Faculty and Student?

Covering constraints Must every superclass entity *also belong* to one of its subclasses?
E.g., must every Person entity be a Student, Faculty, and/or Staff?

Document all constraints.

On good design: Attributes versus Entities

Attributes should represent a *single atomic* value.

Atomic here means a base value without internal structure:

No lists, complex objects,

Entities represent *complex objects*: multiple attributes.

Example: Phone numbers

If you want to model a Person entity with

- ▶ a *single phone number*: use an attribute.
- ▶ *multiple phone numbers*: use a phone number (weak) entity.


On good design: Redundancy

What *redundancy*: storing the same information in multiple ways.

Why Redundancies cause *inconsistencies* during updates.

Redundancies *waste* storage space.

Example of Redundancy



The instructor name is stored *twice*.


On good design: Redundancy

What *redundancy*: storing the same information in multiple ways.

Why Redundancies cause *inconsistencies* during updates.

Redundancies *waste* storage space.

Example of Redundancy (removed)



On good design: Keys & weak entities

- ▶ Real-world entities might not have a clear *key*.
E.g., not everyone has a government-provided SIN.
- ▶ We often create unique internal identifiers for entity sets.
E.g., Mac ID as used within McMaster University.
E.g., sequential counters in many systems.

Weak entities and normal entities

- ▶ Every weak entity can be turned into a normal entity.
- ▶ Use weak entities *only* if there is a clear identifier in the context of an *owner*.
E.g., the *course code* '2DB3' is identifying a course if you know the program.
E.g., the *shirt number* '5' is identifying a player if you know the sport team.

Summary

Entity-relationship modeling

- ▶ yields a *conceptual schema*: high-level description of the data in terms of:
 - ▶ entities, attributes, and relationships,
 - ▶ weak entities, ISA hierarchies, and aggregation,
 - ▶ constraints: keys, relationship participation, ISA hierarchies;
- ▶ depends on a complete *requirements analysis*;
- ▶ is highly subjective: many ways to model the same data.

Later lectures: Translating an ER-diagram into tables.