

SQL: The Structured Query Language

COMPSCI 2DB3: Databases

Jelle Hellings Holly Koponen

Department of Computing and Software
McMaster University



Winter 2024

What is SQL?

Simple answer

A *query language* for interacting with relational databases.

A *high-level programming language* for relational data in which you specify *what* you want, not *how to do it*.

What is SQL?

Simple answer

A *query language* for interacting with relational databases.

A *high-level programming language* for relational data in which you specify *what* you want, not *how to do it*.

Official answer: An ANSI and ISO standard

Information technology–Database languages–SQL

Part 2: Foundation (SQL/Foundation)

What is SQL?

Simple answer

A *query language* for interacting with relational databases.

A *high-level programming language* for relational data in which you specify *what* you want, not *how to do it*.

Official answer: An ANSI and ISO standard

Information technology–Database languages–SQL

Part 2: Foundation (SQL/Foundation) → 1707 pages!

There are eight other parts.

Corrigendum in 2019 with one additional part.

What is SQL?

Simple answer

A *query language* for interacting with relational databases.

A *high-level programming language* for relational data in which you specify *what* you want, not *how to do it*.

Official answer: An ANSI and ISO standard

Information technology–Database languages–SQL

Part 2: Foundation (SQL/Foundation) → 1707 pages!

There are eight other parts.

Corrigendum in 2019 with one additional part.

We will only look at a *very small* part of SQL.

Source: <https://www.iso.org/standard/63555.html>.

Classes of SQL Statements

DML *Data Manipulation Language*: SQL-data and SQL-data change statements.
Create, modify, delete, and inspect stored data (data in tables).

DDL *Data Definition Language*: SQL-schema statements.
Create, modify, delete, and inspect the schema (structure of tables).

SQL-data and SQL-data change statements

We will first look at *data manipulation*.

SQL-data and SQL-data change statements

We will first look at *data manipulation*.

Four main types of SQL-data and SQL-data change statements

INSERT Insert row(s) into a table.

DELETE Delete row(s) from a table.

UPDATE Update row(s) in a table.

SELECT Retrieve data from table(s).

SQL-data and SQL-data change statements

We will first look at *data manipulation*.

Four main types of SQL-data and SQL-data change statements

INSERT Insert row(s) into a table.

DELETE Delete row(s) from a table.

UPDATE Update row(s) in a table.

SELECT Retrieve data from table(s). ← *Our main focus*

Single table: Retrieve a single table

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses;
```

Single table: Retrieve a single table

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses;
```

-- **SELECT** specifies the columns: * is *all columns*.

Single table: Retrieve a single table

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses;
```

-- **SELECT** specifies the columns: * is *all columns*.
-- **FROM** specifies the table(s): **courses**.

Single table: Retrieve a single table

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output		
cid	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses;
```

-- **SELECT** specifies the columns: * is *all columns*.
-- **FROM** specifies the table(s): **courses**.

Returns a copy of the table as-is.

Single table: Retrieve specific columns

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT lecturer  
FROM courses;
```

-- **FROM** specifies the table(s): **courses**.

Single table: Retrieve specific columns

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT lecturer  
FROM courses;
```

```
-- SELECT specifies the columns: lecturer.  
-- FROM specifies the table(s): courses.
```

Single table: Retrieve specific columns

courses				Query output	
<u>cid</u>	title	lecturer		lecturer	
1	Programming	1	→	1	
2	Discrete Mathematics	3		3	
3	Databases	2		2	
4	Advanced Databases	2		2	

SELECT lecturer
FROM courses;

-- **SELECT** specifies the columns: lecturer.
-- **FROM** specifies the table(s): **courses**.

Returns a table with a copy of the column lecturer from courses.

Single table: Retrieve specific columns

courses				Query output	
<u>cid</u>	title	lecturer		lecturer	
1	Programming	1	→	1	} <i>Multiset!</i>
2	Discrete Mathematics	3		3	
3	Databases	2		2	
4	Advanced Databases	2		2	

```
SELECT lecturer  
FROM courses;
```

```
-- SELECT specifies the columns: lecturer.  
-- FROM specifies the table(s): courses.
```

Returns a table with a copy of the column lecturer from courses.

Single table: Retrieve specific columns

courses				
<u>cid</u>	title	lecturer		
1	Programming	1		
2	Discrete Mathematics	3		
3	Databases	2		
4	Advanced Databases	2		

→

Query output	
lecturer	
1	} <i>Set</i>
3	
2	

SELECT DISTINCT lecturer -- **SELECT** specifies the columns: lecturer.
FROM courses; -- **FROM** specifies the table(s): **courses**.

Returns a table with a copy of the column lecturer from courses (unique values).

Single table: Retrieving specific rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses  
WHERE lecturer = 2;
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses.
```

Single table: Retrieving specific rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses  
WHERE lecturer = 2;
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses.  
-- WHERE specifies conditions on each row.
```

Single table: Retrieving specific rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output		
cid	title	lecturer
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses  
WHERE lecturer = 2;
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses.  
-- WHERE specifies conditions on each row.
```

Returns a copy of the rows in the table that meet the condition.

Single table: Retrieving specific rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output		
cid	title	lecturer
2	Discrete Mathematics	3

```
SELECT *  
FROM courses  
WHERE cid < lecturer;
```

- **SELECT** specifies the columns: * is all columns.
- **FROM** specifies the table(s): **courses**.
- **WHERE** specifies conditions on each row.
- Conditions don't need to make sense!

Returns a copy of the rows in the table that meet the condition.

Single table: Retrieving specific rows (text)

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *                -- SELECT specifies the columns: * is all columns.  
FROM courses           -- FROM specifies the table(s): courses.  
WHERE title = 'Databases'; -- WHERE specifies conditions on each row.
```

Single table: Retrieving specific rows (text)

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses  
WHERE title = 'Databases';
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses.  
-- WHERE specifies conditions on each row.  
-- 'Exact' text match.
```


Single table: Retrieving specific rows (text)

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output		
cid	title	lecturer
3	Databases	2

```
SELECT *  
FROM courses  
WHERE title = 'Databases';
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses.  
-- WHERE specifies conditions on each row.  
-- 'Exact' text match.
```

Returns a copy of the rows in the table that meet the text condition.

Single table: Retrieving specific rows (text)

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output		
cid	title	lecturer
3	Databases	2

```
SELECT *  
FROM courses  
WHERE title = 'databases';
```

- **SELECT** specifies the columns: * is all columns.
- **FROM** specifies the table(s): **courses**.
- **WHERE** specifies conditions on each row.
- 'Exact' text match (influenced by the collation).

Returns a copy of the rows in the table that meet the text condition.

Single table: Retrieving specific rows (text)

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output		
cid	title	lecturer
3	Databases	2
4	Advanced Databases	2

```
SELECT *  
FROM courses  
WHERE title LIKE  
        '%Databases%';
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses.  
-- WHERE specifies conditions on each row.  
-- Text containing 'Databases': % matches any text.
```

Returns a copy of the rows in the table that meet the text condition.

Single table: Computations on returned columns

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *, cid + 5,  
        cid * lecturer
```

-- **SELECT** specifies the columns: * is all columns.

```
FROM courses;
```

-- **FROM** specifies the table(s): **courses**.

Single table: Computations on returned columns

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT *, cid + 5,  
        cid * lecturer
```

-- **SELECT** specifies the columns: * is all columns.
-- New columns can be computed out of existing data.

```
FROM courses;
```

-- **FROM** specifies the table(s): **courses**.

Single table: Computations on returned columns

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output				
cid	title	lecturer	<i>C1</i>	<i>C2</i>
1	Programming	1	6	1
2	Discrete Mathematics	3	7	6
3	Databases	2	8	6
4	Advanced Databases	2	9	8

```
SELECT *, cid + 5,  
        cid * lecturer
```

```
FROM courses;
```

-- **SELECT** specifies the columns: * is all columns.
-- New columns can be computed out of existing data.

-- **FROM** specifies the table(s): **courses**.

Returns the table with two added columns computed from the data.

Single table: Computations on returned columns

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output				
cid	title	lecturer	<i>one</i>	<i>two</i>
1	Programming	1	6	1
2	Discrete Mathematics	3	7	6
3	Databases	2	8	6
4	Advanced Databases	2	9	8

```
SELECT *, cid + 5 AS one,  
        cid * lecturer  
        AS two  
FROM courses;
```

```
-- SELECT specifies the columns: * is all columns.  
-- New columns can be computed out of existing data.  
-- AS will rename a column in the output table.  
-- FROM specifies the table(s): courses.
```

Returns the table with two added columns computed from the data (named columns).

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

‘Return courses together with the name of their lecturer.’

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

‘Return courses together with the name of their lecturer.’

```
SELECT *           -- SELECT specifies the columns: * is all columns.
FROM courses, faculty; -- FROM specifies the table(s): courses, faculty.
```

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

'Return courses together with the name of their lecturer.'

SELECT * -- **SELECT** specifies the columns: * is all columns.
FROM courses, faculty; -- **FROM** specifies the table(s): **courses**, **faculty**.

Return the table with all combinations of rows from courses and faculty.

Query output					
cid	title	lecturer	fid	name	rank
2	Discrete Mathematics	3	2	Bo	Assistant
2	Discrete Mathematics	3	3	Celeste	Associate
3	Databases	2	2	Bo	Assistant
3	Databases	2	3	Celeste	Associate

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

'Return courses together with the name of their lecturer.'

SELECT * -- **SELECT** specifies the columns: * is all columns.
FROM courses, faculty; -- **FROM** specifies the table(s): **courses**, **faculty**.

Return the table with all combinations of rows from courses and faculty.

Query output					
cid	title	lecturer	fid	name	rank
2	Discrete Mathematics	3	2	Bo	Assistant
2	Discrete Mathematics	3	3	Celeste	Associate
3	Databases	2	2	Bo	Assistant
3	Databases	2	3	Celeste	Associate

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

‘Return courses together with the name of their lecturer.’

```
SELECT *  
FROM courses, faculty  
WHERE lecturer = fid;
```

```
-- SELECT specifies the columns: * is all columns.  
-- FROM specifies the table(s): courses, faculty.  
-- WHERE specifies conditions on each row.
```

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

'Return courses together with the name of their lecturer.'

SELECT *	-- SELECT specifies the columns: * is all columns.
FROM courses, faculty	-- FROM specifies the table(s): courses , faculty .
WHERE lecturer = fid;	-- WHERE specifies conditions on each row.

Return the table with some combinations of rows from courses and faculty.

Query output					
cid	title	lecturer	fid	name	rank
2	Discrete Mathematics	3	3	Celeste	Associate
3	Databases	2	2	Bo	Assistant

Selecting data from multiple tables

courses		
<u>cid</u>	title	lecturer
2	Discrete Mathematics	3
3	Databases	2

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate

‘Return courses together with the name of their lecturer.’

SELECT title, name -- **SELECT** specifies the columns: title and name only.
FROM courses, faculty -- **FROM** specifies the table(s): **courses**, **faculty**.
WHERE lecturer = fid; -- **WHERE** specifies conditions on each row.

Return the table with some columns from some combinations of rows from courses and faculty.

Query output	
title	name
Discrete Mathematics	Celeste
Databases	Bo

SQL: You specify what you want, not how to do it

```
SELECT title, name  
FROM courses, faculty  
WHERE lecturer = fid;
```

Same query in a structured programming language

```
for  $i = 0$  to |courses| do  
  for  $j = 0$  to |faculty| do  
    if courses[ $i$ ].lecturer = faculty[ $j$ ].fid then  
      output (courses[ $i$ ].title = faculty[ $j$ ].name)
```


SQL: You specify what you want, not how to do it

```
SELECT title, name  
FROM courses, faculty  
WHERE lecturer = fid;
```

Same query in a structured programming language

```
for  $i = 0$  to  $|courses|$  do  
  for  $j = 0$  to  $|faculty|$  do  
    if  $courses[i].lecturer = faculty[j].fid$  then  
      output ( $courses[i].title = faculty[j].name$ )
```

Assumption: *courses* is an in-memory array.

Assumption: *faculty* is an in-memory array.

SQL: You specify what you want, not how to do it

```
SELECT title, name  
FROM courses, faculty  
WHERE lecturer = fid;
```

Same query in a structured programming language

```
for  $i = 0$  to  $|courses|$  do  
  for  $j = 0$  to  $|faculty|$  do  
    if  $courses[i].lecturer = faculty[j].fid$  then  
      output ( $courses[i].title = faculty[j].name$ )
```

} *Costly: $O(|courses| \times |faculty|)$*

SQL: You specify what you want, not how to do it

```
SELECT title, name  
FROM courses, faculty  
WHERE lecturer = fid;
```

Same query in a structured programming language

```
for  $i = 0$  to  $|courses|$  do  
  for  $j = 0$  to  $|faculty|$  do  
    if  $courses[i].lecturer = faculty[j].fid$  then  
      output ( $courses[i].title = faculty[j].name$ )
```

} *Costly: $O(|courses| \times |faculty|)$*

Faster alternative: index structure (hash map, search tree) to look up *faculty*.

SQL: You specify what you want, not how to do it

```
SELECT title, name  
FROM courses, faculty  
WHERE lecturer = fid;
```

Same query in a structured programming language

```
for  $i = 0$  to  $|courses|$  do  
  for  $j = 0$  to  $|faculty|$  do  
    if  $courses[i].lecturer = faculty[j].fid$  then  
      output ( $courses[i].title = faculty[j].name$ )
```

} *Costly: $O(|courses| \times |faculty|)$*

Faster alternative: index structure (hash map, search tree) to look up *faculty*.

Completely different code!

SQL: You specify what you want, not how to do it

```
SELECT title, name  
FROM courses, faculty  
WHERE lecturer = fid;
```

Same query in a structured programming language

```
for  $i = 0$  to  $|courses|$  do  
  for  $j = 0$  to  $|faculty|$  do  
    if  $courses[i].lecturer = faculty[j].fid$  then  
      output ( $courses[i].title = faculty[j].name$ )
```

} *Costly: $O(|courses| \times |faculty|)$*

Faster alternative: index structure (hash map, search tree) to look up *faculty*.

Completely different code!

SQL query: complexity likely very low, e.g., $O(|result| + \log\text{-terms})$.

Sorting returned rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

```
SELECT title, lecturer    -- Default: ordering of output is implementation defined.  
FROM courses;           -- E.g., due to how data is stored and algorithms used.
```

Sorting returned rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output	
title	lecturer
Advanced Databases	2
Databases	2
Discrete Mathematics	3
Programming	1

SELECT title, lecturer -- Default: ordering of output is implementation defined.
FROM courses -- E.g., due to how data is stored and algorithms used.
ORDER BY title; -- You can specify an explicit ordering of the output.

Sorting returned rows

courses		
<u>cid</u>	title	lecturer
1	Programming	1
2	Discrete Mathematics	3
3	Databases	2
4	Advanced Databases	2

→

Query output	
title	lecturer
Discrete Mathematics	3
Advanced Databases	2
Databases	2
Programming	1

```
SELECT title, lecturer  -- Default: ordering of output is implementation defined.
FROM courses           -- E.g., due to how data is stored and algorithms used.
ORDER BY lecturer DESC, title ASC;
                        -- The ordering can be DESCending or ASCending.
                        -- ASCending is the default if not specified.
```


Intermission: Syntax for creating and deleting tables

```
CREATE TABLE table name  
(  
    column specifications & optional constraint specifications  
);
```

Intermission: Syntax for creating and deleting tables

```
CREATE TABLE table name  
(  
    column specifications & optional constraint specifications  
);
```

```
DROP TABLE table name;
```

Intermission: Creating a normal table

```
CREATE TABLE reviews
(  
  pid INT NOT NULL,  
  rid INT NOT NULL,  
  quality INT NOT NULL,  
  originality INT NOT NULL,  
  description CLOB NOT NULL,  
  
  PRIMARY KEY(pid, rid)  
);
```

Intermission: Creating a normal table

```
CREATE TABLE reviews      -- Create a table named 'reviews'.  
(  
  pid INT NOT NULL,  
  rid INT NOT NULL,  
  quality INT NOT NULL,  
  originality INT NOT NULL,  
  description CLOB NOT NULL,  
  
  PRIMARY KEY(pid, rid)  
);
```

Intermission: Creating a normal table

```
CREATE TABLE reviews      -- Create a table named 'reviews'.  
(  
  pid INT NOT NULL,        -- Column 'pid' has type INT.  
  rid INT NOT NULL,  
  quality INT NOT NULL,  
  originality INT NOT NULL,  
  description CLOB NOT NULL,  
  
  PRIMARY KEY(pid, rid)  
);
```

Intermission: Creating a normal table

```
CREATE TABLE reviews      -- Create a table named 'reviews'.
(
  pid INT NOT NULL,          -- Column 'pid' has type INT.
  rid INT NOT NULL,
  quality INT NOT NULL,
  originality INT NOT NULL,
  description CLOB NOT NULL,
                                -- CHARACTER LARGE OBJECT: a lot of text.
  PRIMARY KEY(pid, rid)
);
```

Intermission: Creating a normal table

```
CREATE TABLE reviews      -- Create a table named 'reviews'.  
(  
  pid INT NOT NULL,        -- Column 'pid' has type INT.  
  rid INT NOT NULL,        -- NOT NULL: must have a value.  
  quality INT NOT NULL,  
  originality INT NOT NULL,  
  description CLOB NOT NULL,  
                                -- CHARACTER LARGE OBJECT: a lot of text.  
  PRIMARY KEY(pid, rid)  
);
```

Intermission: Creating a normal table

```
CREATE TABLE reviews      -- Create a table named 'reviews'.  
(  
  pid INT NOT NULL,         -- Column 'pid' has type INT.  
  rid INT NOT NULL,         -- NOT NULL: must have a value.  
  quality INT NOT NULL,  
  originality INT NOT NULL,  
  description CLOB NOT NULL,  
                                -- CHARACTER LARGE OBJECT: a lot of text.  
  PRIMARY KEY(pid, rid)    -- The primary key consists of two columns.  
);
```


Intermission: Creating a normal table

```
CREATE TABLE reviewswn      -- Create a table named 'reviewswn'.  
(  
  pid INT NOT NULL,          -- Column 'pid' has type INT.  
  rid INT NOT NULL,          -- NOT NULL: must have a value.  
  quality INT,              -- Might not have values!  
  originality INT,          -- Might not have values!  
  description CLOB NOT NULL, -- CHARACTER LARGE OBJECT: a lot of text.  
  PRIMARY KEY(pid, rid)      -- The primary key consists of two columns.  
);
```

Intermission: Creating a table with default values

```
CREATE TABLE users          -- Create a table named 'users'.  
(  
  uid INT GENERATED ALWAYS AS IDENTITY,  
  
  name VARCHAR(100) NOT NULL,  
  
  level INT NOT NULL DEFAULT 13,  
  
  mail VARCHAR(100)  
);
```

Intermission: Creating a table with default values

```
CREATE TABLE users          -- Create a table named 'users'.  
(  
  uid INT GENERATED ALWAYS AS IDENTITY,  
                                     -- An identifier is generated automatically.  
  
  name VARCHAR(100) NOT NULL,  
  
  level INT NOT NULL DEFAULT 13,  
  
  mail VARCHAR(100)  
);
```

Intermission: Creating a table with default values

```
CREATE TABLE users          -- Create a table named 'users'.  
(  
  uid INT GENERATED ALWAYS AS IDENTITY,  
                                     -- An identifier is generated automatically.  
  
  name VARCHAR(100) NOT NULL,  
                                     -- A small amount of text (up-to-100 characters).  
  level INT NOT NULL DEFAULT 13,  
  
  mail VARCHAR(100)  
);
```

Intermission: Creating a table with default values

```
CREATE TABLE users          -- Create a table named 'users'.  
(  
  uid INT GENERATED ALWAYS AS IDENTITY,  
                                     -- An identifier is generated automatically.  
  
  name VARCHAR(100) NOT NULL,  
                                     -- A small amount of text (up-to-100 characters).  
  level INT NOT NULL DEFAULT 13,  
                                     -- A default value of 13.  
  mail VARCHAR(100)  
);
```

Intermission: Creating a table with default values

```
CREATE TABLE users          -- Create a table named 'users'.  
(  
  uid INT GENERATED ALWAYS AS IDENTITY,  
                                     -- An identifier is generated automatically.  
  
  name VARCHAR(100) NOT NULL,  
                                     -- A small amount of text (up-to-100 characters).  
  level INT NOT NULL DEFAULT 13,  
                                     -- A default value of 13.  
  mail VARCHAR(100)          -- The default is the value NULL.  
);
```

Intermission: Creating a table with default values

```
CREATE TABLE users          -- Create a table named 'users'.  
(  
  uid INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,  
                                     -- An identifier is generated automatically.  
                                     -- The primary key consists of a single column.  
  name VARCHAR(100) NOT NULL,  
                                     -- A small amount of text (up-to-100 characters).  
  level INT NOT NULL DEFAULT 13,  
                                     -- A default value of 13.  
  mail VARCHAR(100)           -- The default is the value NULL.  
);
```

Intermission: NULL values

Meaning of NULL values

Applications give meaning to NULL values, not SQL. E.g.,

- ▶ optional value or value not applicable;
- ▶ missing value or corrupt value;
- ▶ something else entirely.

SQL has specific rules on how it deals with NULL values in all situations.

NULL values complicate *everything*!

Use them sparingly and prefer NOT NULL columns.

Intermission: Syntax for inserting rows of data

```
INSERT INTO table name  
(column names)optional  
VALUES (values for the row; one value per column)  
    , (more rows) ... optional ;
```

Intermission: Inserting rows of data

reviews

pid rid quality originality description

```
INSERT INTO reviews  
VALUES(1, 1, 6, 7, 'Great!');
```

Intermission: Inserting rows of data

reviews

pid rid quality originality description

```
INSERT INTO reviews      -- Add a row to the table named 'review'.  
VALUES(1, 1, 6, 7, 'Great!'); -- The values for the new row.
```

Intermission: Inserting rows of data

reviews

pid rid quality originality description

```
INSERT INTO reviews          -- Add a row to the table named 'review'.
VALUES(1, 1, 6, 7, 'Great!'),  -- The values for the new row.
      (1, 2, 9, 3, 'Amazing!'), -- You can also add several rows at once.
      (2, 1, 4, 9, 'Perfect!'),
      (2, 3, 10, 10, 'Superb!');
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)  
VALUES('Bo');
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)  
VALUES('Bo');
```

```
-- Add a row to the table named 'users'.  
-- Only provide values for the column 'name'.
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');                   -- Only provide values for the column 'name'.  
  
INSERT INTO users(name, level) VALUES('Celeste', 15);
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');
```

```
INSERT INTO users(name, level) VALUES('Celeste', 15);  
INSERT INTO users(name, level, mail) VALUES('Alicia', 9, 'a@b.com');
```


Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');
```

```
INSERT INTO users(name, level) VALUES('Celeste', 15);  
INSERT INTO users(name, level, mail) VALUES('Alicia', 9, 'a@b.com');  
INSERT INTO users(level, name) VALUES(20, 'Dafni');
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');                    -- Only provide values for the column 'name'.  
  
INSERT INTO users VALUES(12, 'Eva', 13, 'e@b.com');
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');                    -- Only provide values for the column 'name'.  
  
INSERT INTO users VALUES(12, 'Eva', 13, 'e@b.com');  
                                -- Will not work, 'uid' must be generated!
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');                   -- Only provide values for the column 'name'.  
  
INSERT INTO users(level, mail) VALUES(17, 'f@b.com');
```

Intermission: Inserting rows of data with default values

users			
<u>uid</u>	name	level	mail
(generated)	(no default)	(13)	(NULL)

```
INSERT INTO users(name)           -- Add a row to the table named 'users'.  
VALUES('Bo');                    -- Only provide values for the column 'name'.  
  
INSERT INTO users(level, mail) VALUES(17, 'f@b.com');  
                                -- Will not work, 'name' has no default!
```

Intermission: Syntax for deleting rows of data

```
DELETE FROM table name  
WHERE conditions on the rows to deleteoptional ;
```

Intermission: Deleting rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

Intermission: Deleting rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
DELETE FROM reviews WHERE rid = 3;
```


Intermission: Deleting rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
DELETE FROM reviews WHERE pid = rid;
```

Intermission: Deleting rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
DELETE FROM reviews;
```

Intermission: Syntax for updating rows of data

UPDATE *table name*

SET *column name = new value*

more columns to update ... optional

WHERE *conditions on the rows to update*_{optional} ;

Intermission: Updating rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

Intermission: Updating rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	0	10	Superb!

```
UPDATE reviews SET quality = 0 WHERE rid = 3;
```

Intermission: Updating rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	10	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
UPDATE reviews SET quality = quality + 1 WHERE rid = 2;
```

Intermission: Updating rows of data

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	7	6	Great!
1	2	3	9	Amazing!
2	1	9	4	Perfect!
2	3	10	10	Superb!

UPDATE reviews **SET** quality = originality, originality = quality;

This is well-defined in the standard—some DBMSs (MySQL) are broken and do not support this.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	2
2	3

```
SELECT pid, rid
FROM reviews
WHERE quality > 8;
```

Return review-rows with a high (≥ 8) quality.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	2
2	3

```
SELECT pid, rid
```

```
FROM reviews
```

```
WHERE quality > 8;
```

```
-- Alternatives: quality >= 9, NOT(quality <= 8), and NOT quality < 9.
```

Return review-rows with a high (≥ 8) quality.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
2	3

```
SELECT pid, rid
FROM reviews
WHERE quality > 8 AND originality > 8;
```

Return review-rows with a high (≥ 8) quality and originality.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!



Query output	
<u>pid</u>	<u>rid</u>
1	2
2	1
2	3

```
SELECT pid, rid
FROM reviews
WHERE quality > 8 OR originality > 8;
```

Return review-rows with a high (≥ 8) quality or originality.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	1

```
SELECT pid, rid
FROM reviews
WHERE NOT(quality > 8 OR originality > 8);
```

Return review-rows without a high (≥ 8) quality or originality.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	1

```
SELECT pid, rid
```

```
FROM reviews
```

```
WHERE NOT(quality > 8 OR originality > 8);
```

```
-- Alternative: NOT(quality > 8) AND NOT(originality > 8).
```

Return review-rows without a high (≥ 8) quality or originality.

A closer look at comparisons in **WHERE**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	1
1	2
2	1

```
SELECT pid, rid
FROM reviews
WHERE quality <> 10;
```

Return review-rows without the highest quality.

Queries and NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid, quality + originality AS x  
FROM reviewswn  
WHERE quality > 8 OR originality > 8;
```

Question: Which pairs (pid, rid) do you expect in the output?

Vote at <https://strawpoll.com/pdaw6rza6>.

Or: go to <https://strawpoll.live> and use the code **268762**.

Comparisons with NULL values

Question: What should the result of `NULL = NULL` and of `NULL <> NULL` be?

Vote at <https://strawpoll.com/w3hycpxzq>.

Or: go to <https://strawpoll.live> and use the code **733519**.

Comparisons with NULL values

Question: What should the result of `NULL = NULL` and of `NULL <> NULL` be?

Vote at <https://strawpoll.com/w3hycpxzq>.

Or: go to <https://strawpoll.live> and use the code **733519**.

- ▶ SQL does not know what a NULL value means.
- ▶ Comparing with a NULL value has an *unknown* result.
- ▶ SQL does not know whether that *unknown* is true or false.

All comparisons with NULL result in unknown!

Comparisons with NULL values

Question: What should the result of `NULL = NULL` and of `NULL <> NULL` be?

Vote at <https://strawpoll.com/w3hycpxzq>.

Or: go to <https://strawpoll.live> and use the code **733519**.

- ▶ SQL does not know what a NULL value means.
- ▶ Comparing with a NULL value has an *unknown* result.
- ▶ SQL does not know whether that *unknown* is true or false.

All comparisons with NULL result in unknown!

DISTINCT *does* consider different NULL values to be equal.

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

Rules for **AND**

- ▶ $A \text{ AND true} \equiv A$.
- ▶ $A \text{ AND false} \equiv \text{false}$.
- ▶ $\text{unknown AND unknown} \equiv \text{unknown}$.

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

Rules for **OR**

- ▶ $A \text{ OR true} \equiv \text{true}$.
- ▶ $A \text{ OR false} \equiv A$.
- ▶ $\text{unknown OR unknown} \equiv \text{unknown}$.

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

Rules for **NOT**

- ▶ **NOT** true = false.
- ▶ **NOT** false = true.
- ▶ **NOT** unknown = unknown.

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

A truth table for the three-valued logic of SQL

<i>A</i>	<i>B</i>	<i>A AND B</i>	<i>A OR B</i>	NOT <i>A</i>
false	false	false	false	true
true	false	false	true	false
false	true	false	true	true
true	true	true	true	false

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

A truth table for the three-valued logic of SQL

<i>A</i>	<i>B</i>	<i>A AND B</i>	<i>A OR B</i>	NOT <i>A</i>
false	false	false	false	true
true	false	false	true	false
false	true	false	true	true
true	true	true	true	false
false	unknown	false	unknown	true
true	unknown	unknown	true	false

The three-valued (ternary) logic of SQL

A logic with three values: true, false, and unknown.

A truth table for the three-valued logic of SQL

<i>A</i>	<i>B</i>	<i>A AND B</i>	<i>A OR B</i>	NOT <i>A</i>
false	false	false	false	true
true	false	false	true	false
false	true	false	true	true
true	true	true	true	false
false	unknown	false	unknown	true
true	unknown	unknown	true	false
unknown	false	false	unknown	unknown
unknown	true	unknown	true	unknown
unknown	unknown	unknown	unknown	unknown

Computations with NULL

- ▶ Simple computations: if the input has NULL, then the output is NULL.
E.g., $5 + \text{NULL} = \text{NULL}$, $\text{NULL} * 8 = \text{NULL}$, ...
- ▶ Aggregate operations (which we cover later): NULL is *mostly* ignored.

Answering queries with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid, quality + originality AS x
FROM reviewswn
WHERE quality > 8 OR originality > 8;
```

Answering queries with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid, quality + originality AS x
FROM reviewswn
WHERE quality > 8 OR originality > 8;
```

Answering queries with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid, quality + originality AS x
FROM reviewswn
WHERE quality > 8 OR originality > 8;
```

Answering queries with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

→

Query output		
pid	rid	x
2	1	NULL
2	3	NULL

```
SELECT pid, rid, quality + originality AS x
FROM reviewswn
WHERE quality > 8 OR originality > 8;
```

Working with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid
FROM reviewswn
WHERE quality IS NULL AND
      originality IS NOT NULL;
```

Working with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid
FROM reviewswn
WHERE quality IS NULL AND -- Check whether 'quality' is a NULL value.
      originality IS NOT NULL;
```


Working with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

```
SELECT pid, rid
FROM reviewswn
WHERE quality IS NULL AND -- Check whether 'quality' is a NULL value.
      originality IS NOT NULL;
                        -- Check whether 'originality' is not a NULL value.
```

Working with NULL values

reviewswn				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	NULL	NULL	Amazing!
2	1	NULL	9	Perfect!
2	3	10	NULL	Superb!

→

Query output	
pid	rid
2	1

```
SELECT pid, rid
```

```
FROM reviewswn
```

```
WHERE quality IS NULL AND -- Check whether 'quality' is a NULL value.  
      originality IS NOT NULL;
```

```
-- Check whether 'originality' is not a NULL value.
```

Set-wise combining **SELECT** queries

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

```
SELECT sid AS id FROM students;
```

Return the identifier of all students.

Set-wise combining **SELECT** queries

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

```
SELECT sid AS id FROM students
UNION
SELECT fid AS id FROM faculty;
```

Return the identifier of all students and all faculty.

Set-wise combining **SELECT** queries

students			faculty			→	Query output
<u>sid</u>	name	year	<u>fid</u>	name	rank		id
1	Alicia	2020	2	Bo	Assistant		1
3	Celeste	2018	3	Celeste	Associate		2
4	Dafni	2019	5	Eva	Full		3
							4
							5

SELECT sid **AS** id **FROM** students

UNION

-- **UNION** combines results as a *set*.

SELECT fid **AS** id **FROM** faculty;

Return the identifier of all students and all faculty.

Set-wise combining **SELECT** queries

students			faculty			→	Query output
<u>sid</u>	name	year	<u>fid</u>	name	rank		id
1	Alicia	2020	2	Bo	Assistant		1
3	Celeste	2018	3	Celeste	Associate		2
4	Dafni	2019	5	Eva	Full		3
							3
							4
							5

SELECT sid **AS** id **FROM** students

UNION ALL

-- **UNION ALL** combines results as a *multiset*.

SELECT fid **AS** id **FROM** faculty;

Return the identifier of all students and all faculty.

Set-wise combining **SELECT** queries

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

→

Query output	
<u>id</u>	
3	

```
SELECT sid AS id FROM students
```

```
INTERSECT
```

-- **INTERSECT** combines results as a *set*.

```
SELECT fid AS id FROM faculty;
```

Return the identifier of all students that are faculty.

Set-wise combining **SELECT** queries

students			faculty			→	Query output	
<u>sid</u>	name	year	<u>fid</u>	name	rank		id	
1	Alicia	2020	2	Bo	Assistant		1	
3	Celeste	2018	3	Celeste	Associate		4	
4	Dafni	2019	5	Eva	Full			

SELECT sid **AS** id **FROM** students

EXCEPT -- **EXCEPT** combines results as a *set*.

SELECT fid **AS** id **FROM** faculty;

Return the identifier of all students that are not faculty.

The set and multiset semantics of SQL

A	B
v	v
1	1
2	2
2	2
3	3
3	3
3	3
4	3
4	4
4	4

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.

The set and multiset semantics of SQL

A		B	
v		v	
1		1	
2		2	
2		2	
3		3	
3		3	
3		3	
4		3	
4		4	
4		4	

UNION

→

Query output
v
1
2
3
4

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.
- ▶ **UNION:** W will be once in the output if W is in the input (if $m + n > 0$).

The set and multiset semantics of SQL

		Query output
A	B	v
v	v	
1	1	1
2	2	2
2	2	2
3	3	2
3	3	3
3	3	3
4	3	3
4	4	3
4	4	3
		3
		3
		...

UNION ALL

→

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.
- ▶ **UNION ALL:** W will be once in the output for every copy of W in the input ($m + n$ output copies).

The set and multiset semantics of SQL

A		B	
v		v	
1	INTERSECT	1	Query output v
2		2	
2		2	1
3		3	
3		3	
3		3	
4		3	2
4		4	
4		4	
4		4	3
			4

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.
- ▶ **INTERSECT:** W will be once in the output if it is in both inputs (if $\min(m, n) > 0$).

The set and multiset semantics of SQL

A		B		Query output
v		v		v
1	INTERSECT ALL	1	→	1
2		2		2
2		2		2
3		3		3
3		3		3
3		3		3
4		3		4
4		4		4
4		4		4
4		4		4

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.
- ▶ **INTERSECT ALL:** W will be once in the output for every copy of W in both inputs ($\min(m, n)$ output copies).

The set and multiset semantics of SQL

A		B	
v		v	
1		1	
2		2	
2		2	
3	EXCEPT	3	
3		3	
3		3	
4		3	
4		4	
4		4	

→

Query output
v

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.
- ▶ **EXCEPT:** W will be once in the output if it is only in the left input (if $m > 0 = n$).

The set and multiset semantics of SQL

A		B	
v		v	
1		1	
2		2	
2		2	
3	EXCEPT ALL	3	
3		3	
3		3	
4		3	
4		4	
4		4	

→

Query output
v
4

Consider a value W

- ▶ Let m be the count of W in 'A'.
- ▶ Let n be the count of W in 'B'.
- ▶ **EXCEPT ALL:** W will be once in the output for every extra copy in the left input ($m - n$ copies if $m > n$).

The set and multiset semantics of SQL

A
v
1
2
2
3
3
3
4
4
4

B
v
1
2
2
3
3
3
3
4
4

What about NULL values?

The set and multiset semantics of SQL

A
v
1
2
2
3
3
3
4
4
4

B
v
1
2
2
3
3
3
3
4
4

What about NULL values?
Are you doing the right thing?

The set and multiset semantics of SQL

A
v
1
2
2
3
3
3
4
4
4

B
v
1
2
2
3
3
3
3
4
4

What about NULL values?

Are you doing the right thing?

NULLs are treated as normal values.

Nested queries: Another option for intersections and differences

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

```
SELECT sid
FROM students
WHERE sid IN (                -- IN checks whether 'sid' is in a list
    SELECT fid FROM faculty); -- obtained from column 'fid' in table 'faculty'.
```

Nested queries: Another option for intersections and differences

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

→

Query output
id
3

```
SELECT sid
FROM students
WHERE sid IN (
    SELECT fid FROM faculty);
```

-- **IN** checks whether 'sid' is in a list
-- obtained from column 'fid' in table 'faculty'.

Return the identifier of all students that are faculty.

Nested queries: Another option for intersections and differences

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

Query output
id
1
4

```
SELECT sid
FROM students
WHERE sid NOT IN (                -- NOT IN checks whether 'sid' is not in a list
  SELECT fid FROM faculty);      -- obtained from column 'fid' in table 'faculty'.
```

Return the identifier of all students that are not faculty.

Nested queries: Another option for intersections and differences

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

Query output
id
1
4


```
SELECT sid
FROM students
WHERE sid NOT IN (                -- NOT IN checks whether 'sid' is not in a list
  SELECT fid FROM faculty);      -- obtained from column 'fid' in table 'faculty'.
```

Return the identifier of all students that are not faculty.

Nested queries are power tools that, with proper usage, are highly efficient!

Nesting nested queries

friendof
<u>from</u> <u>to</u>




```
SELECT "FROM" FROM friendof  
WHERE to = 'Frieda';
```

- "FROM" is the column name 'from' and is not
- interpreted as the **FROM** keyword.
- Double-quoted names are *case sensitive*!
- DB2 stores names upper-case internally.


Nesting nested queries

friendof
<u>from</u> <u>to</u>




```
SELECT "FROM" FROM friendof  
WHERE to IN (  
  SELECT "FROM" FROM friendof  
  WHERE to = 'Frieda');
```


Nesting nested queries



```
SELECT "FROM" FROM friendof
WHERE to IN (
  SELECT "FROM" FROM friendof
  WHERE to IN (
    SELECT "FROM" FROM friendof
    WHERE to = 'Frieda'));
```

Nesting nested queries



```
SELECT "FROM" FROM friendof
WHERE to IN (
  SELECT "FROM" FROM friendof
  WHERE to IN (
    SELECT "FROM" FROM friendof
    WHERE to IN (
      SELECT "FROM" FROM friendof
      WHERE to = 'Frieda'))));
```

Nested queries and **ALL**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
SELECT pid, rid
FROM reviews
WHERE originality >= ALL (
    SELECT originality      -- The nested query gives a list L of
    FROM reviews);         -- all 'originality' scores.
```

Nested queries and **ALL**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
SELECT pid, rid
FROM reviews
WHERE originality >= ALL ( -- >= ALL: 'originality' must be max(L).
    SELECT originality    -- The nested query gives a list L of
    FROM reviews);       -- all 'originality' scores.
```

Nested queries and **ALL**

reviews					→	Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description		<u>pid</u>	<u>rid</u>
1	1	6	7	Great!			
1	2	9	3	Amazing!			
2	1	4	9	Perfect!			
2	3	10	10	Superb!			

```
SELECT pid, rid
FROM reviews
WHERE originality >= ALL ( -- >= ALL: 'originality' must be max(L).
    SELECT originality    -- The nested query gives a list L of
    FROM reviews);       -- all 'originality' scores.
```

Return the primary keys of review(s) with the highest originality score.

Nested queries and **ALL**

reviews					→	Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description		<u>pid</u>	<u>rid</u>
1	1	6	7	Great!		1	2
1	2	9	3	Amazing!			
2	1	4	9	Perfect!			
2	3	10	10	Superb!			

```
SELECT pid, rid
```

```
FROM reviews
```

```
WHERE originality < ALL ( -- < ALL: 'originality' must be less than min(L).  
    SELECT quality        -- The nested query gives a list L of  
    FROM reviews);       -- all 'quality' scores.
```

Return the primary keys of review(s) with an originality score lower than any quality score.

Nested queries and **ALL**

reviews					→	Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description		<u>pid</u>	<u>rid</u>
1	1	6	7	Great!		1	1
1	2	9	3	Amazing!		1	2
2	1	4	9	Perfect!			
2	3	10	10	Superb!			

```
SELECT pid, rid
FROM reviews
WHERE originality <> ALL ( -- <> ALL: 'originality' must be NOT IN L.
    SELECT quality        -- The nested query gives a list L of
    FROM reviews);        -- all 'quality' scores.
```

Return the primary keys of review(s) with an originality score that is not a quality score.

Correlated nested queries

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
SELECT pid, rid
FROM reviews
WHERE originality >= ALL (
    SELECT originality
    FROM reviews);
```

Return, for each reviewer, the primary keys of its review(s) with the highest originality score.

Correlated nested queries

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
2	3

```
SELECT pid, rid
FROM reviews
WHERE originality >= ALL (
    SELECT originality
    FROM reviews);
```

Return, for each reviewer, the primary keys of its review(s) with the highest originality score.

Correlated nested queries

reviews					→	Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description		<u>pid</u>	<u>rid</u>
1	1	6	7	Great!			
1	2	9	3	Amazing!			
2	1	4	9	Perfect!			
2	3	10	10	Superb!			

```
SELECT pid, rid
FROM reviews AS ra          -- Give the two 'reviews' different names.
WHERE originality >= ALL (
    SELECT originality
    FROM reviews AS rb);    -- Give the two 'reviews' different names.
```

Return, for each reviewer, the primary keys of its review(s) with the highest originality score.

Correlated nested queries

reviews					→	Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description		<u>pid</u>	<u>rid</u>
1	1	6	7	Great!			
1	2	9	3	Amazing!			
2	1	4	9	Perfect!			
2	3	10	10	Superb!			

```
SELECT pid, rid
FROM reviews ra           -- Give the two 'reviews' different names.
WHERE originality >= ALL ( -- AS is always optional.
    SELECT originality
    FROM reviews rb);     -- Give the two 'reviews' different names.
```

Return, for each reviewer, the primary keys of its review(s) with the highest originality score.

Correlated nested queries

reviews					→	Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description		<u>pid</u>	<u>rid</u>
1	1	6	7	Great!		1	2
1	2	9	3	Amazing!		2	1
2	1	4	9	Perfect!		2	3
2	3	10	10	Superb!			


```
SELECT pid, rid
FROM reviews AS ra          -- Give the two 'reviews' different names.
WHERE originality >= ALL (   -- AS is always optional.
    SELECT originality
    FROM reviews AS rb      -- Give the two 'reviews' different names.
    WHERE ra.rid = rb.rid); -- Only 'originality' scores of this reviewer.
```

Return, for each reviewer, the primary keys of its review(s) with the highest originality score.

Nesting correlated nested queries

friendof
<u>from</u> <u>to</u>


```
SELECT "FROM" FROM friendof F
WHERE to = "FROM";
```



Nesting correlated nested queries


friendof
<u>from</u> <u>to</u>

```
SELECT "FROM" FROM friendof F
WHERE to IN (
  SELECT "FROM" FROM friendof
  WHERE to = F."FROM");
```



Nesting correlated nested queries


friendof
<u>from</u> <u>to</u>



```
SELECT "FROM" FROM friendof F
WHERE to IN (
  SELECT "FROM" FROM friendof
  WHERE to = F."FROM"));
```

Nesting correlated nested queries


friendof
<u>from</u> <u>to</u>



```
SELECT "FROM" FROM friendof F
WHERE to IN (
  SELECT "FROM" FROM friendof
  WHERE to IN (
    SELECT "FROM" FROM friendof
    WHERE to IN (
      SELECT "FROM" FROM friendof
      WHERE to = F."FROM"))));
```


Nesting correlated nested queries

friendof
<u>from</u> <u>to</u>




```
SELECT "FROM" FROM friendof F
WHERE to IN (
  SELECT "FROM" FROM friendof
  WHERE to IN (
    SELECT "FROM" FROM friendof
    WHERE to IN (
      SELECT "FROM" FROM friendof
      WHERE to = F."FROM"))));
```

How to do a 4-length Hamiltonian cycle?

Nesting correlated nested queries

friendof
<u>from</u> <u>to</u>



```
SELECT "FROM" FROM friendof F
WHERE to IN (
  SELECT "FROM" FROM friendof F2
  WHERE to IN (
    SELECT "FROM" FROM friendof F3
    WHERE to IN (
      SELECT "FROM" FROM friendof F4
      WHERE to = F."FROM" AND
        F."FROM" <> F2."FROM" AND
        F."FROM" <> F3."FROM" AND
        F."FROM" <> F4."FROM" AND
        F2."FROM" <> F3."FROM" AND
        F2."FROM" <> F4."FROM" AND
        F3."FROM" <> F4."FROM"))));
```


Nesting correlated nested queries

friendof
<u>from</u> <u>to</u>

How to do a 4-length Hamiltonian cycle?

We also still have normal joins!

```
SELECT F."FROM", F2."FROM",  
       F3."FROM", F4."FROM"  
FROM friendof F, friendof F2,  
     friendof F3, friendof F4  
WHERE F.to = F2."FROM" AND F2.to = F3."FROM"  
     AND F3.to = F4."FROM" AND F4.to = F."FROM"  
     AND F."FROM" <> F2."FROM" AND F."FROM" <> F3."FROM"  
     AND F."FROM" <> F4."FROM" AND F2."FROM" <> F3."FROM"  
     AND F2."FROM" <> F4."FROM" AND F3."FROM" <> F4."FROM";
```




Return the nodes on a 4-length Hamiltonian cycle .

Nesting correlated nested queries

friendof
<u>from</u> <u>to</u>

How to do a 4-length Hamiltonian cycle?



We also still have normal joins!

```
SELECT F."FROM", F2."FROM",  
       F3."FROM", F4."FROM"  
FROM friendof F, friendof F2,  
     friendof F3, friendof F4  
WHERE F.to = F2."FROM" AND F2.to = F3."FROM"  
     AND F3.to = F4."FROM" AND F4.to = F."FROM"  
     AND F."FROM" <> F2."FROM" AND F."FROM" <> F3."FROM"  
     AND F."FROM" <> F4."FROM" AND F2."FROM" <> F3."FROM"  
     AND F2."FROM" <> F4."FROM" AND F3."FROM" <> F4."FROM"  
     AND F."FROM" < F2."FROM" AND F."FROM" < F3."FROM"  
     AND F."FROM" < F4."FROM";
```

Return the nodes on a 4-length Hamiltonian cycle once.

Nested queries and **ANY**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
2	1
2	3

```
SELECT pid, rid
```

```
FROM reviews
```

```
WHERE originality = ANY ( -- = ANY: 'originality' must be IN L.
```

```
    SELECT quality        -- The nested query gives a list L of
```

```
    FROM reviews);       -- all 'quality' scores.
```

Return the review(s) with an originality score equal to a quality score.

Nested queries and **ANY**

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	1
2	1
2	3

```
SELECT pid, rid
```

```
FROM reviews
```

```
WHERE originality > ANY ( -- > ANY: 'originality' must be more than min(L).  
    SELECT quality        -- The nested query gives a list L of  
    FROM reviews);       -- all 'quality' scores.
```

Return the review(s) with an originality score larger than the minimum quality score.

Nested queries and **ANY**

reviews					Query output	
<u>pid</u>	<u>rid</u>	quality	originality	description	<u>pid</u>	<u>rid</u>
1	1	6	7	Great!	1	1
1	2	9	3	Amazing!	1	2
2	1	4	9	Perfect!	2	1
2	3	10	10	Superb!	2	3

```
SELECT pid, rid
FROM reviews
WHERE originality <> ANY (
    SELECT quality
    FROM reviews);
```

-- <> **ANY**: a different value must be in *L*.
-- The nested query gives a list *L* of
-- all 'quality' scores.

Return the review(s) with an originality score different from a quality score.

Nested queries and **EXIST**

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

```
SELECT * FROM students  
WHERE EXISTS (  
  SELECT *  
  FROM faculty);
```

-- The nested query gives a list L with
-- all rows in 'faculty'.

Nested queries and **EXIST**

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

→

Query output		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

```
SELECT * FROM students
WHERE EXISTS (
  SELECT *
  FROM faculty);
```

- **EXISTS**: true if $L \neq \emptyset$.
- The nested query gives a list L with
- all rows in 'faculty'.

Return the students if there are faculty members.

Nested queries and **EXIST**

students		
<u>sid</u>	name	year
1	Alicia	2020
3	Celeste	2018
4	Dafni	2019

faculty		
<u>fid</u>	name	rank
2	Bo	Assistant
3	Celeste	Associate
5	Eva	Full

→

Query output		
<u>sid</u>	name	year
1	Alicia	2020
4	Dafni	2019

```
SELECT * FROM students
WHERE NOT EXISTS (
  SELECT *
  FROM faculty
  WHERE sid = fid);
```

- **NOT EXISTS**: true if $L = \emptyset$.
- The nested query gives a list L with
- all rows in 'faculty' that are the same
- person as the student.

Return those students that are not also faculty members.

BETWEEN: A feature that is easily forgotten

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

```
SELECT pid, rid
FROM review
WHERE originality BETWEEN 7 AND 9;
```

BETWEEN: A feature that is easily forgotten

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	1
2	1

```
SELECT pid, rid
FROM review
WHERE originality BETWEEN 7 AND 9;
```

Returns those reviews with an originality score between 7 and 9.

BETWEEN: A feature that is easily forgotten

reviews				
<u>pid</u>	<u>rid</u>	quality	originality	description
1	1	6	7	Great!
1	2	9	3	Amazing!
2	1	4	9	Perfect!
2	3	10	10	Superb!

→

Query output	
<u>pid</u>	<u>rid</u>
1	1
2	1

```
SELECT pid, rid
```

```
FROM review
```

```
WHERE originality BETWEEN 7 AND 9;
```

```
-- Equivalent to 7 <= originality AND originality <= 9.
```

Returns those reviews with an originality score between 7 and 9.

Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

```
SELECT *  
FROM courses C, instructors I  
WHERE C.cid = I.cid;
```


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
2	D. Mathematics	2	Eva
3	Databases	3	Alicia



```
SELECT *  
FROM courses C, instructors I  
WHERE C.cid = I.cid;
```

-- Output only contains those courses and lectures
-- that are connected to each other.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
2	D. Mathematics	2	Eva
3	Databases	3	Alicia



```
SELECT *  
FROM courses C CROSS JOIN  
      instructors I  
WHERE C.cid = I.cid;
```

-- Output only contains those courses and lectures
-- that are connected to each other.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
2	D. Mathematics	2	Eva
3	Databases	3	Alicia



```
SELECT *  
FROM courses C JOIN  
      instructors I ON C.cid = I.cid;
```

-- Output only contains those courses and lectures
-- that are connected to each other.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
2	D. Mathematics	2	Eva
3	Databases	3	Alicia



```
SELECT *  
FROM courses C INNER JOIN  
instructors I ON C.cid = I.cid;
```

-- Output only contains those courses and lectures
-- that are connected to each other.

Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

```
SELECT *  
FROM courses C LEFT JOIN  
         instructors I ON C.cid = I.cid;
```


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
1	Programming	NULL	NULL
2	D. Mathematics	2	Eva
3	Databases	3	Alicia



```
SELECT *  
FROM courses C LEFT JOIN  
      instructors I ON C.cid = I.cid;
```

- Output contains *all* courses.
- Courses with instructors include instructor details.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
1	Programming	NULL	NULL
2	D. Mathematics	2	Eva
3	Databases	3	Alicia



```
SELECT *  
FROM courses C LEFT OUTER JOIN  
      instructors I ON C.cid = I.cid;
```

- Output contains *all* courses.
- Courses with instructors include instructor details.

Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

```
SELECT *  
FROM courses C RIGHT JOIN  
        instructors I ON C.cid = I.cid;
```

- Output contains *all* instructors.
- Instructors with courses include course details.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
2	D. Mathematics	2	Eva
3	Databases	3	Alicia
NULL	NULL	4	Bo



```
SELECT *  
FROM courses C RIGHT JOIN  
instructors I ON C.cid = I.cid;
```

- Output contains *all* instructors.
- Instructors with courses include course details.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo



Query output			
cid	title	cid	name
2	D. Mathematics	2	Eva
3	Databases	3	Alicia
NULL	NULL	4	Bo



```
SELECT *  
FROM courses C RIGHT OUTER JOIN  
instructors I ON C.cid = I.cid;
```

- Output contains *all* instructors.
- Instructors with courses include course details.

Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

```
SELECT *  
FROM courses C FULL JOIN  
        instructors I ON C.cid = I.cid;
```

- Output contains *all* courses and instructors.
- Courses and instructors are connected if possible.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

→

Query output			
cid	title	cid	name
1	Programming	NULL	NULL
2	D. Mathematics	2	Eva
3	Databases	3	Alicia
NULL	NULL	4	Bo



```
SELECT *  
FROM courses C FULL JOIN  
      instructors I ON C.cid = I.cid;
```

- Output contains *all* courses and instructors.
- Courses and instructors are connected if possible.


Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

→

Query output			
cid	title	cid	name
1	Programming	NULL	NULL
2	D. Mathematics	2	Eva
3	Databases	3	Alicia
NULL	NULL	4	Bo



```
SELECT *  
FROM courses C FULL OUTER JOIN  
instructors I ON C.cid = I.cid;
```

- Output contains *all* courses and instructors.
- Courses and instructors are connected if possible.

Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

Reminder: Queries don't have to make sense!

SELECT *

FROM courses C **FULL JOIN**

instructors I **ON** C.cid > I.cid;

Selecting data from multiple tables: Revisited

courses	
<u>cid</u>	title
1	Programming
2	D. Mathematics
3	Databases

instructors	
<u>cid</u>	name
2	Eva
3	Alicia
4	Bo

→

Query output			
cid	title	cid	name
1	Programming	NULL	NULL
2	D. Mathematics	NULL	NULL
3	Databases	2	Eva
NULL	NULL	3	Alicia
NULL	NULL	4	Bo

Reminder: Queries don't have to make sense!

SELECT *

FROM courses C **FULL JOIN**

instructors I **ON** C.cid > I.cid;

Aggregation with nested queries

Consider the following two queries

- ▶ Per product: get the review with highest rating.
- ▶ Per user: get the product they rated lowest.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Aggregation with nested queries

Consider the following two queries

- ▶ Per product: get the review with highest rating.
- ▶ Per user: get the product they rated lowest.

Using nested queries

```
SELECT *  
FROM productreview P  
WHERE rating >= ALL (  
    SELECT rating  
    FROM productreview Q  
    WHERE P.product = Q.product);
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Aggregation with nested queries

Consider the following two queries

- ▶ Per product: get the review with highest rating.
- ▶ Per user: get the product they rated lowest.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Using nested queries

```
SELECT *  
FROM productreview P  
WHERE rating <= ALL (  
    SELECT rating  
    FROM productreview Q  
    WHERE P.user = Q.user);
```


Aggregation with nested queries

Consider the following two queries

- ▶ Per product: get the review with highest rating.
- ▶ Per user: get the product they rated lowest.

Using nested queries

Works—but limited and finicky.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

The basics of aggregation

```
SELECT aggregate  
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

The basics of aggregation

```
SELECT MIN(rating)
FROM productreview;
```

-- Minimum value in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Minimum of ratings: 3.

The basics of aggregation

```
SELECT MIN(rating)
FROM productreview;
```

- Minimum value in column 'rating'.
- Aggregates operate on *groups* of rows.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Minimum of ratings: 3.

The basics of aggregation

```
SELECT MIN(rating), product
FROM productreview;
```

- Minimum value in column 'rating'.
- Aggregates operate on *groups* of rows.
- No longer access to individual unaggregated columns!

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Invalid query.

The basics of aggregation

```
SELECT MAX(rating)
FROM productreview;
```

-- Maximum value in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Maximum of ratings: 10.

The basics of aggregation

```
SELECT COUNT(rating)
FROM productreview;
```

-- Number of rows.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Ratings: 7.

The basics of aggregation

```
SELECT COUNT(*)  
FROM productreview;
```

-- Number of rows.

-- For **COUNT**ing rows: the columns don't matter.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Rows: 7.

The basics of aggregation

```
SELECT COUNT(DISTINCT rating)
FROM productreview;
```

-- Number of distinct values in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Distinct ratings: 6.

The basics of aggregation

```
SELECT SUM(rating)  
FROM productreview;
```

-- Sum of the values in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Sum of ratings: 47.

The basics of aggregation

```
SELECT SUM(DISTINCT rating)  
FROM productreview;
```

-- Sum of the distinct values in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Sum of distinct ratings: 42.

The basics of aggregation

```
SELECT AVG(rating)
FROM productreview;
```

-- Average of the values in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Average of ratings: 6.

The basics of aggregation

```
SELECT AVG(rating)
FROM productreview;
```

-- Average of the values in column 'rating'.
-- The average can be an **INT**!

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Average of ratings: 6 (**INT**).

The basics of aggregation

```
SELECT AVG(CAST(rating AS DOUBLE))  
FROM productreview;
```

- Average of the values in column 'rating'.
- The average can be an **INT**!
- Solution: make column 'rating' a **DOUBLE**.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Average of ratings: 6.7 (**DOUBLE**).

The basics of aggregation

```
SELECT AVG(CAST(rating AS DECIMAL))  
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

-- Average of the values in column 'rating'.

-- The average can be an **INT**!

-- Floating point numbers (**DOUBLE**) cause issues.

-- Solution: make column 'rating' a **DECIMAL**.

} Average of ratings: 6.7 (**DECIMAL**).

The basics of aggregation

```
SELECT AVG(DISTINCT rating)
FROM productreview;
```

-- Average of the distinct values in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Average of distinct ratings: 7.

The basics of aggregation

```
SELECT AVG(DISTINCT CAST(rating AS DECIMAL))  
FROM productreview;
```

-- Average of the distinct values in column 'rating'.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

} Average of distinct ratings: 7.0 (**DECIMAL**).

Using basic aggregations in queries

Consider again the following query

Per product: get the review with highest rating.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Using nested queries

```
SELECT *  
FROM productreview P  
WHERE rating >= ALL (  
    SELECT rating  
    FROM productreview Q  
    WHERE P.product = Q.product);
```

Using basic aggregations in queries

Consider again the following query

Per product: get the review with highest rating.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Using aggregation and nested queries

```
SELECT *  
FROM productreview P  
WHERE rating = (  
    SELECT MAX(rating)  
    FROM productreview Q  
    WHERE P.product = Q.product);
```

Using basic aggregations in queries

Consider again the following query

Per product: get the review with highest rating.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Using aggregation and nested queries

```
SELECT *  
FROM productreview P  
WHERE rating = (  
    SELECT MAX(rating)  
    FROM productreview Q  
    WHERE P.product = Q.product);
```

Still a bit underwhelming!

GROUP BY: Putting aggregation to good use

Reminder

Aggregates operate on *groups* of rows.

You can no longer access individual unaggregated columns!

GROUP BY: Putting aggregation to good use

Reminder

Aggregates operate on *groups* of rows \longleftarrow *by default: entire table is one group.*
You can no longer access individual unaggregated columns!

GROUP BY: Putting aggregation to good use

Reminder

Aggregates operate on *groups* of rows \longleftarrow *by default: entire table is one group.*
You can no longer access individual unaggregated columns!

We can specify our own row grouping

```
SELECT aggregated output columns  
FROM sources  
WHERE row conditions  
GROUP BY columns to group rows on;
```

Using **GROUP BY** in queries

Consider the following query

Get the highest and lowest rating per product.

```
SELECT *  
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Using **GROUP BY** in queries

Consider the following query

Get the highest and lowest rating per product.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

```
SELECT *  
FROM productreview  
GROUP BY product;
```

-- We want aggregated information *per product*.

-- Hence, **GROUP BY** product!

Using **GROUP BY** in queries

Consider the following query

Get the highest and lowest rating per product.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
GROUP BY product;
```

-- We want aggregated information *per product*.

-- Hence, **GROUP BY** product!

-- The column '*product*' can be used:

-- every row in *a single aggregated group* is the same product.

Using **GROUP BY** in queries

Consider the following query

Get the highest and lowest rating per product.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
GROUP BY product;
```

Query output		
product	<i>C1</i>	<i>C2</i>
cheese	10	7
phone	5	3
shoe	8	5

Using **GROUP BY** in queries

Consider the following query

Get the highest and lowest rating per product.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
WHERE rating <> 5 -- Filters rows before grouping!
GROUP BY product;
```

Query output		
product	<i>C1</i>	<i>C2</i>
cheese	10	7
phone	<i>3</i>	3
shoe	8	<i>8</i>

HAVING: Filtering on groups

Structure of aggregated queries

```
SELECT aggregated output columns  
FROM sources  
WHERE row conditions  
GROUP BY columns to group rows on;
```

-- Filters rows *before* grouping!

HAVING: Filtering on groups

Structure of aggregated queries

```
SELECT aggregated output columns  
FROM sources  
WHERE row conditions  
GROUP BY columns to group rows on  
HAVING conditions on aggregated groups;
```

-- Filters rows *before* grouping!

HAVING: Filtering on groups

Structure of aggregated queries

SELECT *aggregated output columns*

FROM *sources*

WHERE *row conditions*

-- Filters rows *before* grouping!

GROUP BY *columns to group rows on*

HAVING *conditions on aggregated groups;*

-- Filters aggregated rows

-- (*after* grouping)!

Using **HAVING** in queries

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
GROUP BY product;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Query output		
product	<i>C1</i>	<i>C2</i>
cheese	10	7
phone	5	3
shoe	8	5

Using **HAVING** in queries

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
GROUP BY product
HAVING COUNT(*) > 2;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Using **HAVING** in queries

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
GROUP BY product
HAVING COUNT(*) > 2;
```

-- Only look at products with *many* reviews!

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Query output		
product	<i>C1</i>	<i>C2</i>
cheese	10	7

Using **HAVING** in queries

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
GROUP BY product
HAVING AVG(rating) >= 5;
```

-- Only look at products with *high* average ratings!

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Query output		
product	<i>C1</i>	<i>C2</i>
cheese	10	7
shoe	8	5

Using **HAVING** in queries

```
SELECT product, MAX(rating), MIN(rating)
FROM productreview
WHERE rating > 4
GROUP BY product
HAVING AVG(rating) >= 5;
```

-- First ignore the bad ratings (**WHERE**).
-- Then look at products with high ratings (**HAVING**).

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

Query output		
product	<i>C1</i>	<i>C2</i>
cheese	10	7
phone	5	5
shoe	8	5

The basics of aggregation with NULL

What about NULL values?
Are you doing the right thing?

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

The basics of aggregation with NULL

What about NULL values?

```
SELECT MIN(rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Minimum of ratings: 5.

The basics of aggregation with NULL

What about NULL values?

```
SELECT MAX(rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Maximum of ratings: 10.

The basics of aggregation with NULL

What about NULL values?

```
SELECT COUNT(rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Ratings: 5.

The basics of aggregation with NULL

What about NULL values?

```
SELECT COUNT(*)  
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Rows: 7.

The basics of aggregation with NULL

What about NULL values?

```
SELECT COUNT(DISTINCT rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Distinct ratings: 4.

The basics of aggregation with NULL

What about NULL values?

```
SELECT SUM(rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Sum of ratings: 35.

The basics of aggregation with NULL

What about NULL values?

```
SELECT SUM(DISTINCT rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Sum of distinct ratings: 30.

The basics of aggregation with NULL

What about NULL values?

```
SELECT AVG(rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Average of ratings: 7.

The basics of aggregation with NULL

What about NULL values?

```
SELECT AVG(DISTINCT rating)
FROM productreview;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

} Average of distinct ratings: 7.

The basics of aggregation with NULL

What about NULL values?

```
SELECT MIN(rating)
FROM productreview
WHERE rating IS NULL;
```

-- General rule: NULL values are discarded.
-- Unless all values are NULL!

} Minimum of ratings: NULL.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

The basics of aggregation with NULL

What about NULL values?

```
SELECT MAX(rating)
FROM productreview
WHERE rating IS NULL;
```

-- General rule: NULL values are discarded.
-- Unless all values are NULL!

} Maximum of ratings: NULL.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

The basics of aggregation with NULL

What about NULL values?

```
SELECT COUNT(rating)
FROM productreview
WHERE rating IS NULL;
```

-- General rule: NULL values are discarded.
-- Unless all values are NULL!

Ratings: 0.

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

The basics of aggregation with NULL

What about NULL values?

```
SELECT COUNT(*)  
FROM productreview  
WHERE rating IS NULL;
```

-- General rule: NULL values are discarded.
-- Unless all values are NULL!

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

} Rows: 2.

The basics of aggregation with NULL

What about NULL values?

```
SELECT COUNT(DISTINCT rating)
FROM productreview
WHERE rating IS NULL;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

-- Unless all values are NULL!

} Distinct ratings: 0.

The basics of aggregation with NULL

What about NULL values?

```
SELECT SUM(rating)
FROM productreview
WHERE rating IS NULL;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

-- Unless all values are NULL!

} Sum of ratings: NULL.

The basics of aggregation with NULL

What about NULL values?

```
SELECT SUM(DISTINCT rating)
FROM productreview
WHERE rating IS NULL;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

-- Unless all values are NULL!

} Sum of distinct ratings: NULL.

The basics of aggregation with NULL

What about NULL values?

```
SELECT AVG(rating)
FROM productreview
WHERE rating IS NULL;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

-- Unless all values are NULL!

} Average of ratings: NULL.

The basics of aggregation with NULL

What about NULL values?

```
SELECT AVG(DISTINCT rating)
FROM productreview
WHERE rating IS NULL;
```

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	NULL
Eva	shoe	8
Bo	phone	NULL
Bo	shoe	5
Celeste	cheese	7

-- General rule: NULL values are discarded.

-- Unless all values are NULL!

} Average of distinct ratings: NULL.

The **FROM** clause

The basic **SELECT-FROM-WHERE** structure

SELECT *output columns*

FROM *(joined) source tables*

WHERE *row conditions;*

The **FROM** clause

The basic **SELECT-FROM-WHERE** structure

SELECT *output columns*

FROM *(joined) source tables*

WHERE *row conditions;*

-- Query results can be used as tables!

The **FROM** clause

The basic **SELECT-FROM-WHERE** structure

```
SELECT output columns  
FROM (joined) source tables  
WHERE row conditions;
```

Example

```
SELECT *  
FROM (SELECT * FROM courses  
      WHERE cid IN (SELECT cid FROM instructors))  
WHERE title LIKE '%Databases%';
```

The **FROM** clause

The basic **SELECT-FROM-WHERE** structure

```
SELECT output columns  
FROM (joined) source tables  
WHERE row conditions;
```

Example: Very useful for complex queries with aggregation!

```
SELECT *  
FROM (SELECT product, COUNT(*) AS c, SUM(rating) AS s  
      FROM productreview  
      GROUP BY product) P, sellers S  
WHERE P.product = S.product;
```

The **FROM** clause

The basic **SELECT-FROM-WHERE** structure

```
SELECT output columns  
FROM (joined) source tables  
WHERE row conditions;
```

Example: Very useful for complex queries with aggregation!

```
SELECT *  
FROM (SELECT product, COUNT(*) AS c, SUM(rating) AS s  
      FROM productreview  
      GROUP BY product) P, sellers S  
WHERE P.product = S.product;
```

-- Aggregation *before* the join, otherwise the computed aggregates don't make any sense!

Dealing with complex queries

Question: How do I write *complex* queries?

SQL might look weird, but is still a *programming language*!

Dealing with complex queries

Question: How do I write *complex* queries?

SQL might look weird, but is still a *programming language*!

You write complex programs by *first* breaking them up into subproblems and *then* solving these subproblems independently (e.g., functions).

Dealing with complex queries

Question: How do I write *complex* queries?

SQL might look weird, but is still a *programming language*!

You write complex programs by *first* breaking them up into subproblems and *then* solving these subproblems independently (e.g., functions).

```
WITH name (column names)optional AS (  
    query  
)  
SELECT * FROM name;
```

Dealing with complex queries

Question: How do I write *complex* queries?

SQL might look weird, but is still a *programming language*!

You write complex programs by *first* breaking them up into subproblems and *then* solving these subproblems independently (e.g., functions).

```
WITH name1 (column names)optional AS (  
    query  
) , name2 (column names)optional AS (  
    query  
)  
SELECT * FROM name1 , name2;
```


Dealing with complex queries

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

```
WITH sumresult AS (  
  SELECT product, SUM(rating) AS srating  
  FROM productreview GROUP BY product)  
SELECT * FROM sumresult  
WHERE srating >= ALL (SELECT srating FROM sumresult);
```

Dealing with complex queries

productreview		
<u>user</u>	<u>product</u>	<u>rating</u>
Alicia	cheese	10
Alicia	phone	5
Eva	cheese	9
Eva	shoe	8
Bo	phone	3
Bo	shoe	5
Celeste	cheese	7

→

Query output	
product	srating
cheese	26

```
WITH sumresult AS (  
  SELECT product, SUM(rating) AS srating  
  FROM productreview GROUP BY product)  
SELECT * FROM sumresult  
WHERE srating >= ALL (SELECT srating FROM sumresult);
```

Addendum: Retrieving specific rows based on text patterns

Standard: **SIMILAR TO**

SELECT *

FROM courses

WHERE title **LIKE** '%Databases%';

Addendum: Retrieving specific rows based on text patterns

Standard: **SIMILAR TO**

SELECT *

FROM courses

WHERE title **SIMILAR TO** '%(data|algo)%';

Addendum: Retrieving specific rows based on text patterns

Standard: **SIMILAR TO**

SELECT *

FROM courses

WHERE title **SIMILAR TO** '%(data|algo)%';

- ▶ Standards-compliant support seems to be mostly absent, except for PostgreSQL.
- ▶ Other vendors have vendor-specific options.
(e.g., extensions to **LIKE**, **REGEXP_LIKE**, **REGEXP**, **RLIKE**).

Addendum: Retrieving specific rows based on text patterns

Standard: **SIMILAR TO**

SELECT *

FROM courses

WHERE title **SIMILAR TO** '%(data|algo)%';

- ▶ In general: relational databases are not optimized for *text search*:
searching patterns in unstructured text is the opposite of structured (tabular) data!
- ▶ Searching a few small fields: probably fine (options for index support).
- ▶ Anything beyond: dedicated full-text search and indexing engines are a safer bet.
- ▶ Regular expressions are not a full replacement of proper input checking!
Expressions for names, e-mail, phone numbers, ... are often oversimplified or wrong.